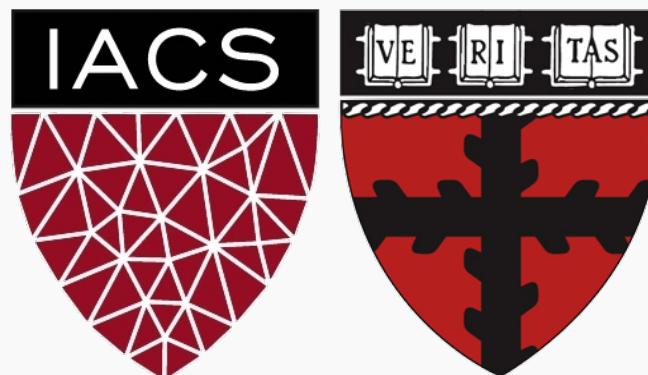


Lecture 29 : Introduction to Generative Adversarial Networks (GANS)

CS109B Data Science 2

Pavlos Protopapas, Mark Glickman, and Chris Tanner





"Generative Adversarial Networks is the **most interesting idea in the last ten years** in machine learning."

Yann LeCun, Director, Facebook AI

Outline

- Generative Modeling Motivation
- High Level Formalism
- Architecture
- Mathematics
- Training GANS
- Conditional GANS



Outline

- **Generative Modeling Motivation**
- High Level Formalism
- Architecture
- Mathematics
- Training GANS
- Conditional GANS

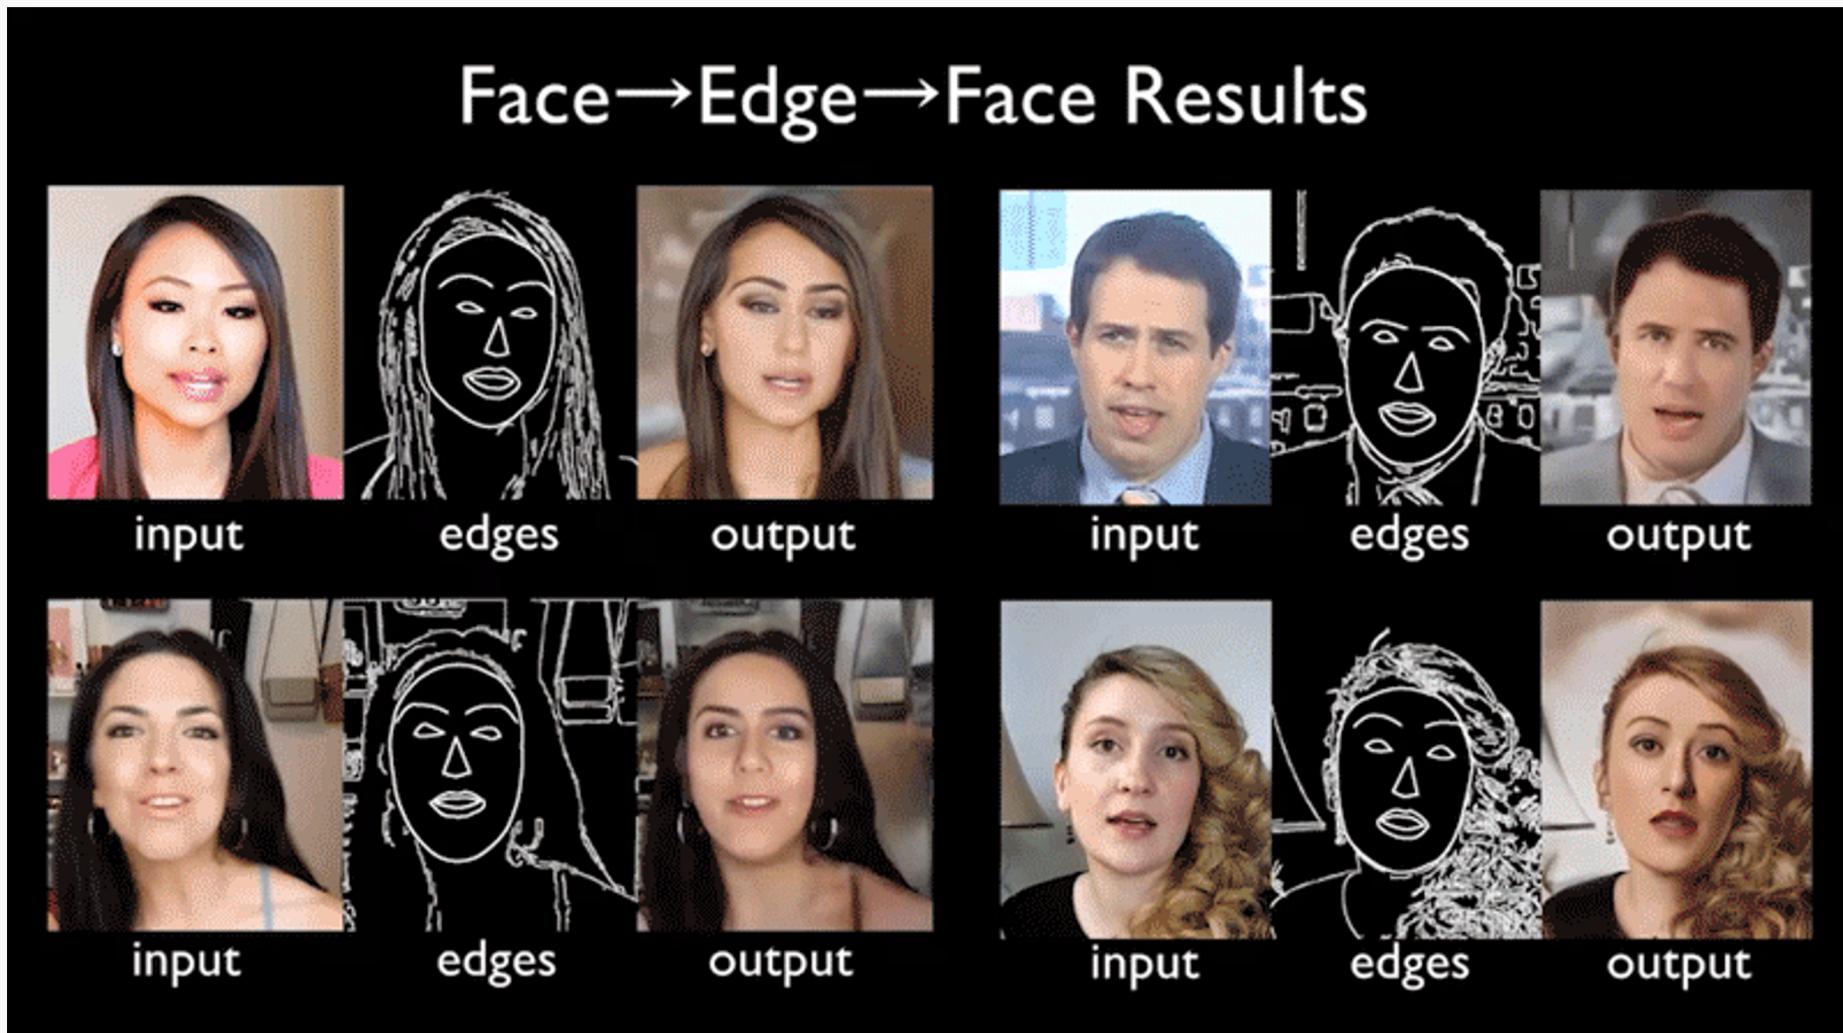


Unpaired Image-to-Image Translation using Cycle-GANs



[\[Zhu et al. 2017\]](#)

Video-to-Video Synthesis



[Wang et al. 2018]

What is generative modeling?

Given samples $\sim p_{\text{data}}$, we would like to sample from the same distribution?



Training data $\sim p_{\text{data}}(x)$



Generated samples $\sim p_{\text{model}}(x)$

What is generative modeling?

How do we generate samples from the same distribution as $p_{\text{data}}(x)$?

Explicit sampling: $p_{\text{model}}(x)$ has analytical expression:

- MCMC
- Variational methods

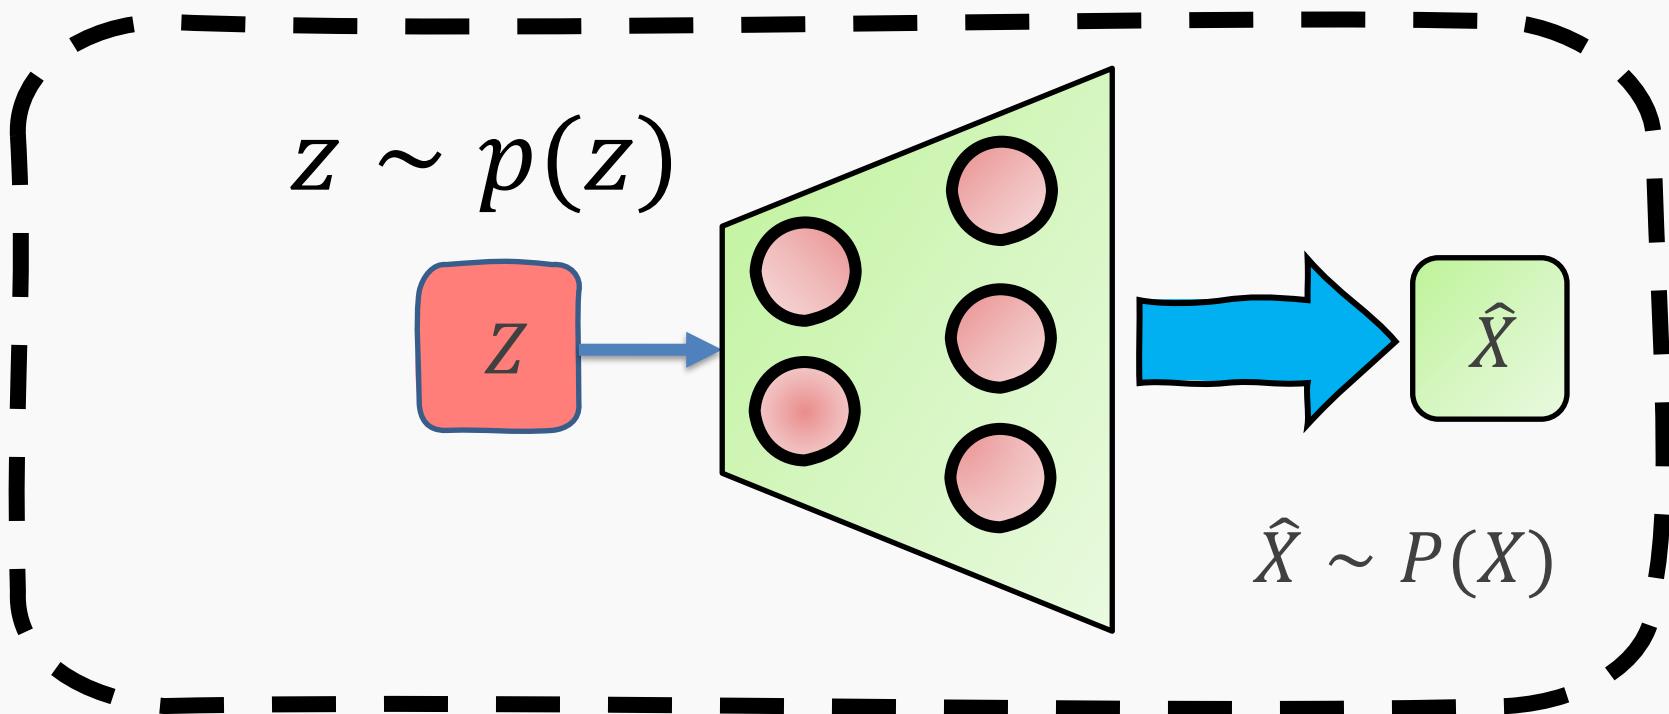
Implicit sampling: learn how to sample from $p_{\text{data}}(x)$ but we do not need the analytical expression:

- Generator part of VAE
- GANS



Variational Auto Encoder as a generative model

Generative model



Though we used Variational inference to sample of the latent space, at the end we created a model that given z in generates \hat{X} with a distribution similar to X .

Why should we study it?

1. Realistic generation tasks
2. Debiasing and data augmentation
3. Missing data
4. Simulation and planning (RL)

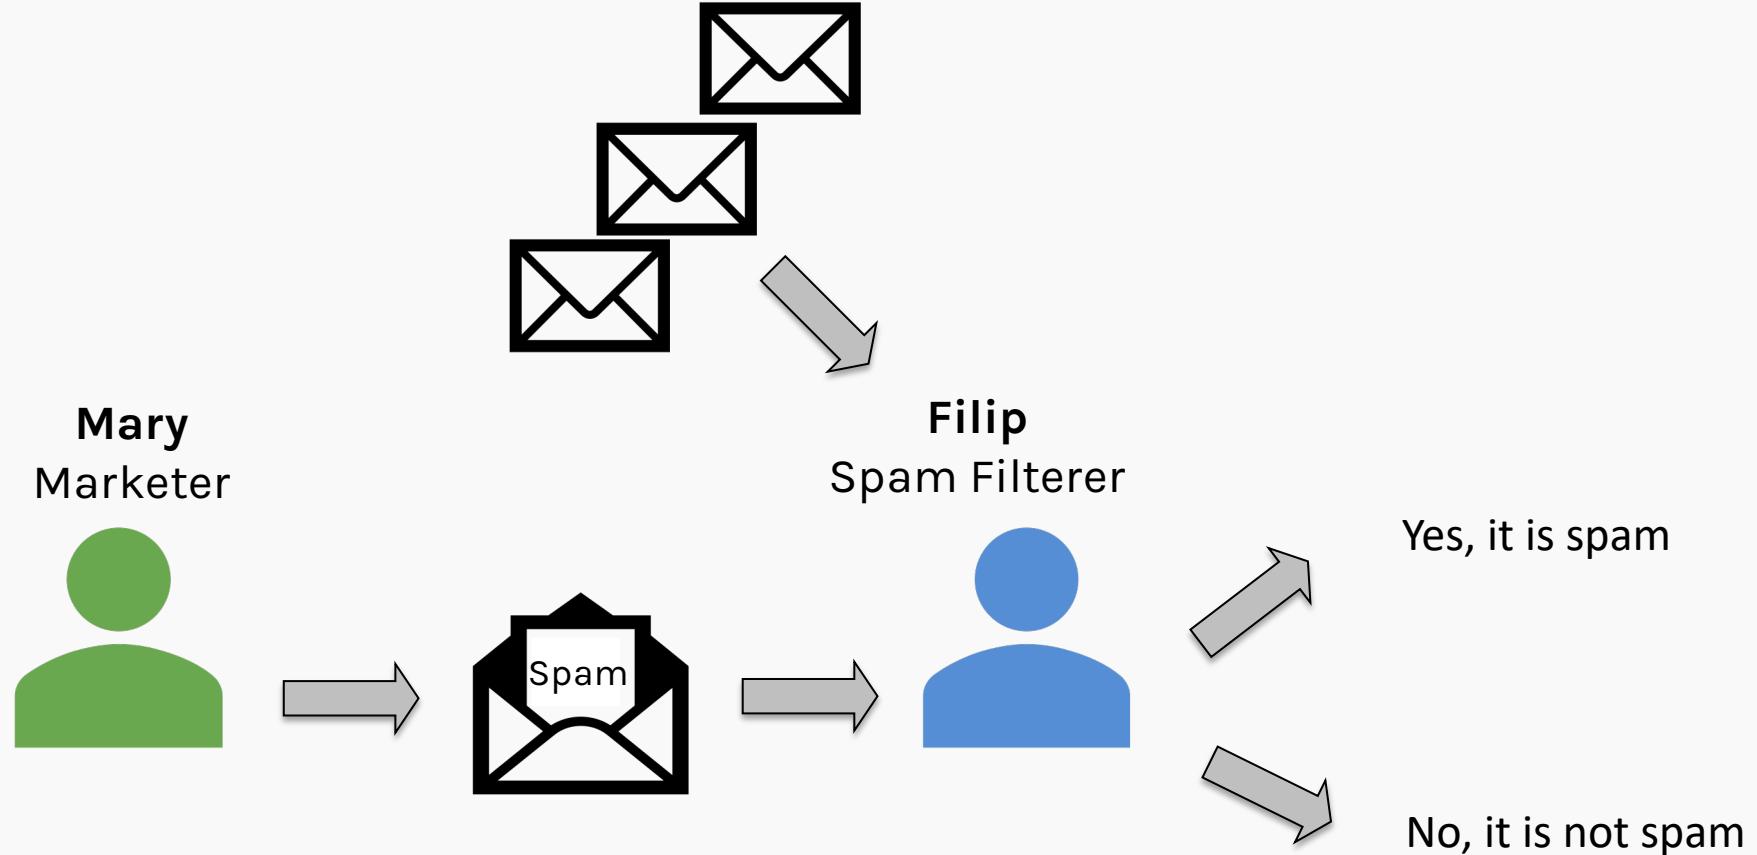


Outline

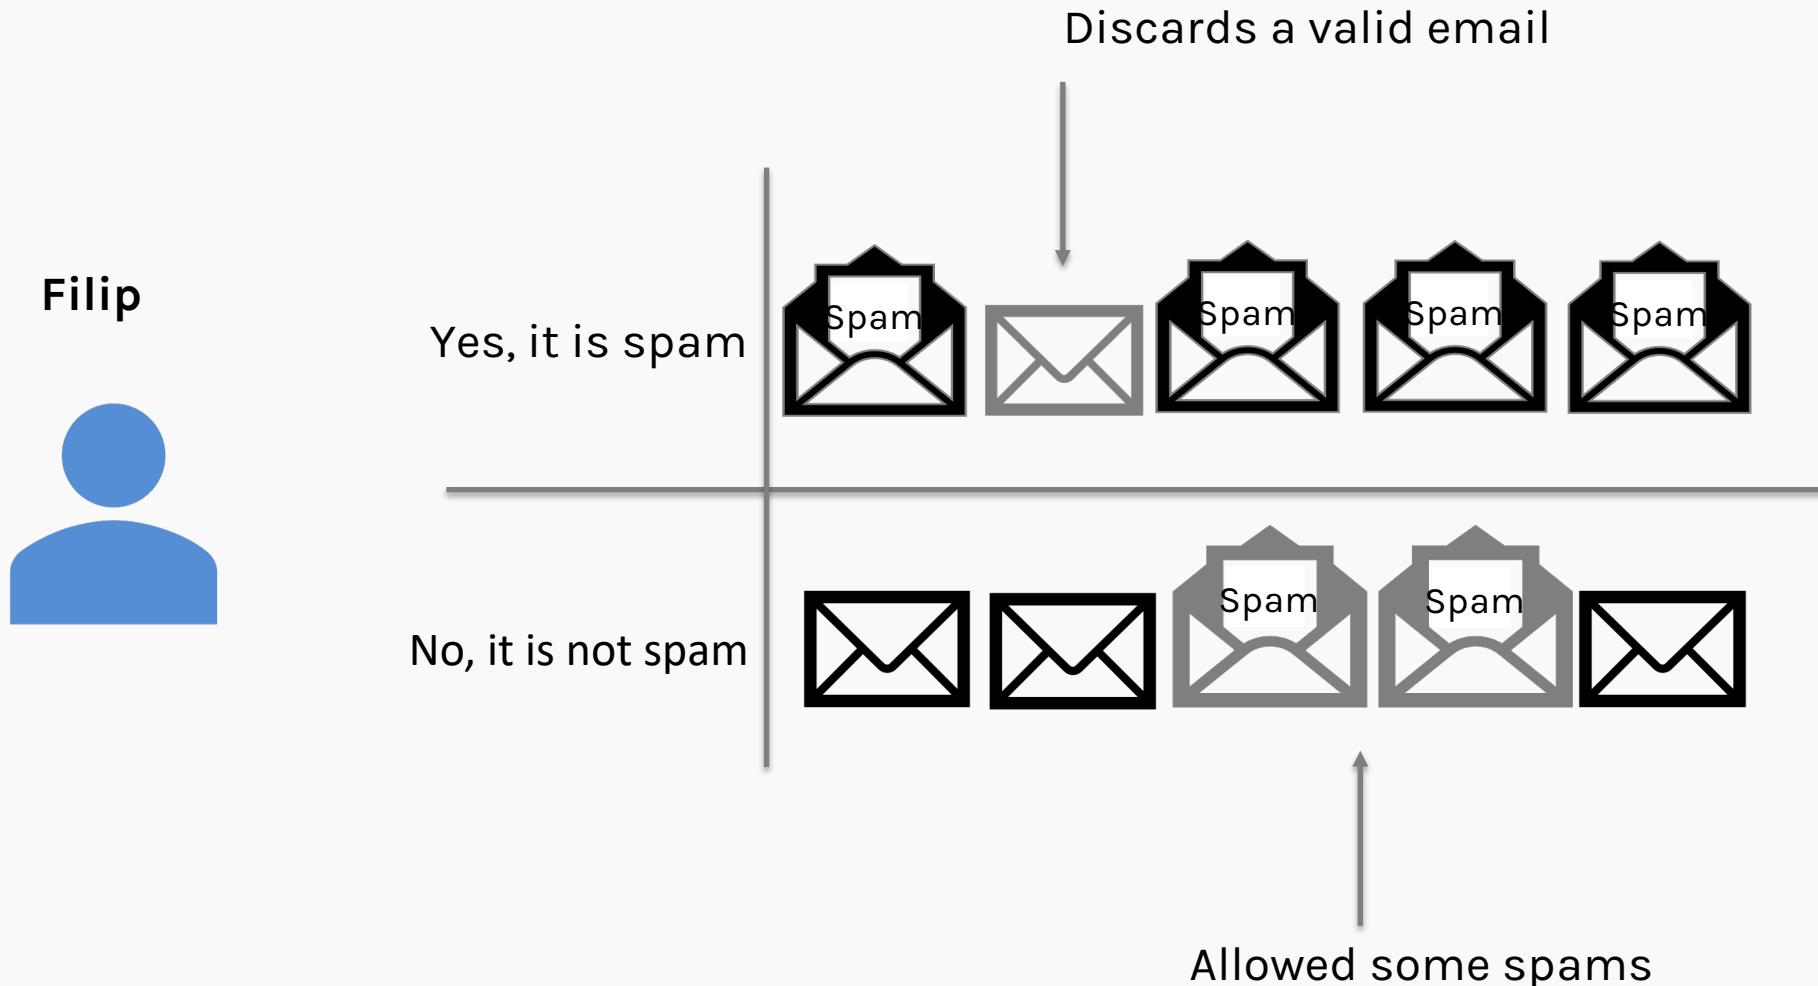
- Generative Modeling Motivation
- **High Level Formalism**
- Architecture
- Mathematics
- Training GANS
- **Conditional GANS**



Generative Adversarial Networks (GANs)



Generative Adversarial Networks (GANs)



Generative Adversarial Networks (GANs)

Mary and Filip
**learned from what
went wrong from
their perspective**



	It was spam, for real	It was not spam
Filip: it is spam	Four black envelope icons with the word "Spam" written inside them, set against a light green background.	One black envelope icon.
Filip: it is not spam	Two grey envelope icons with the word "Spam" written inside them.	Three black envelope icons.

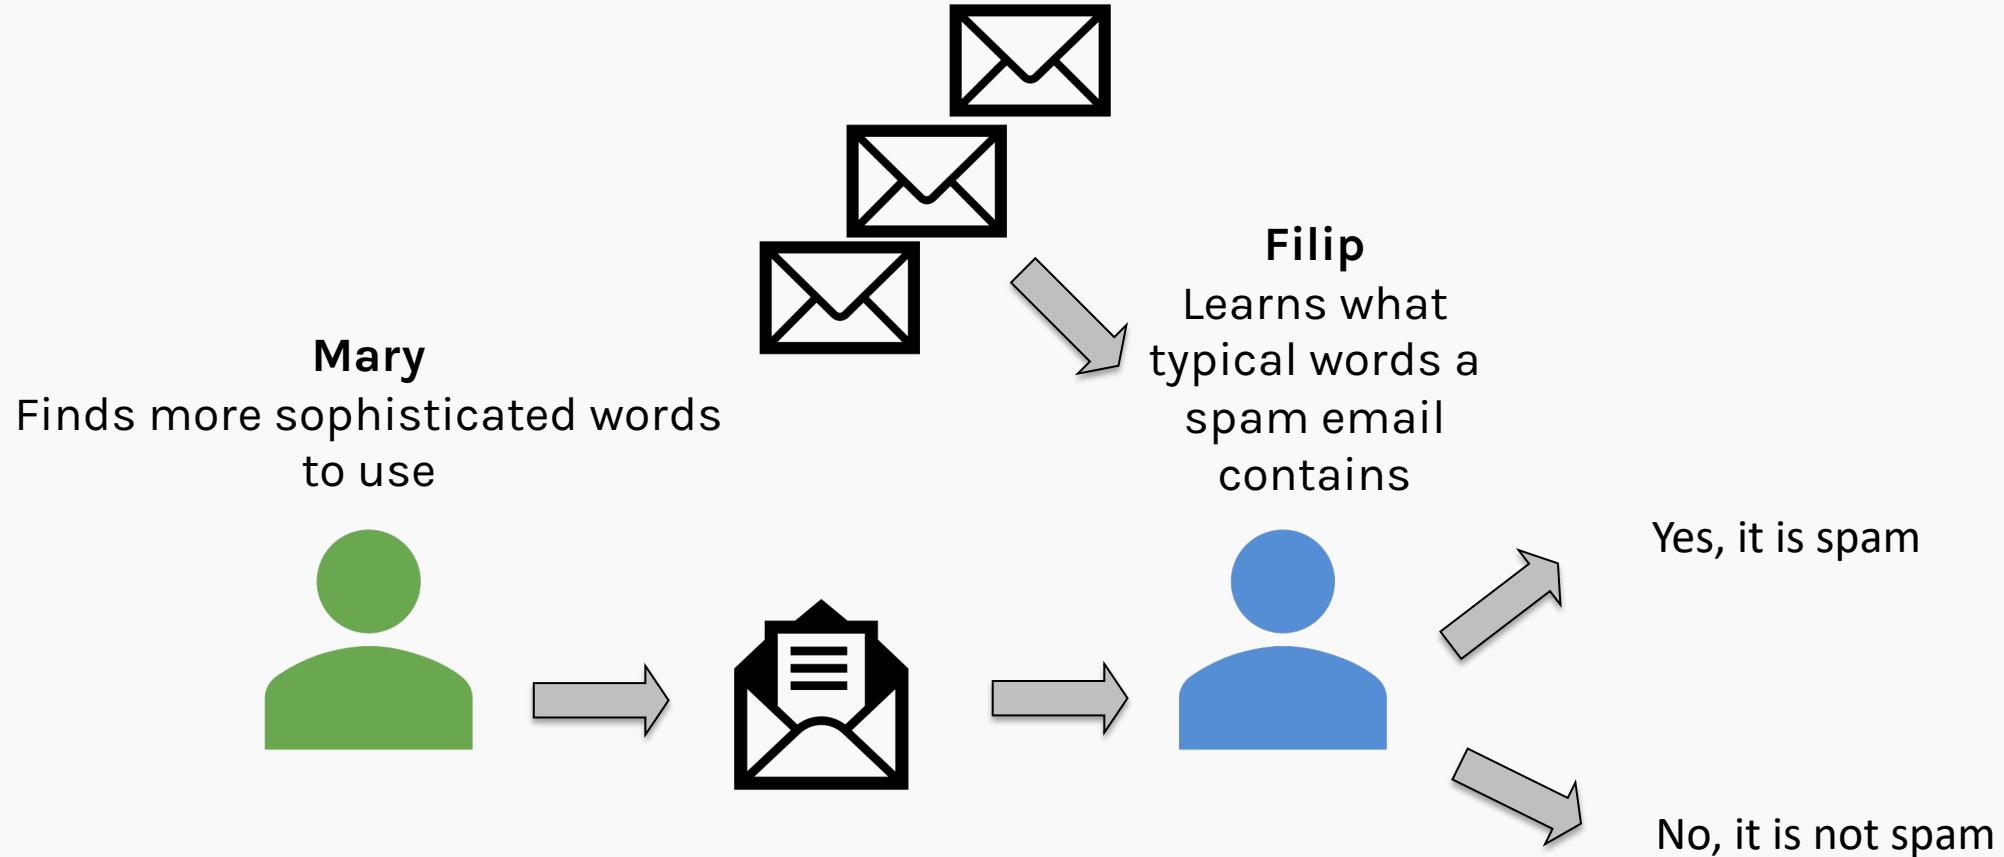
Generative Adversarial Networks (GANs)

Mary and Filip
**learned from what
went wrong from
their perspective**

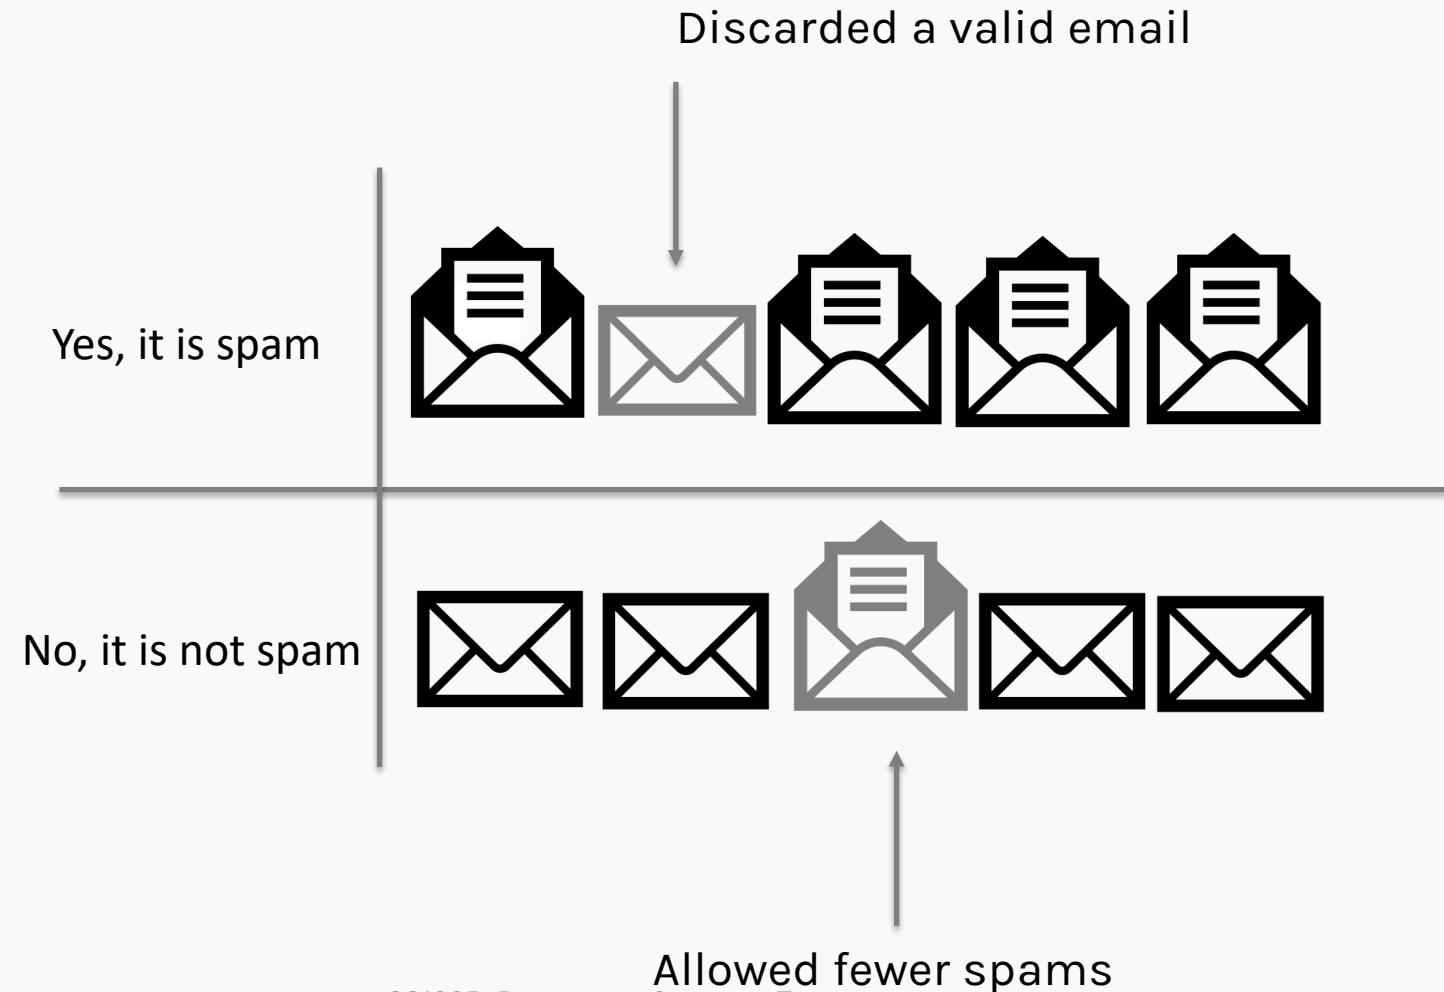
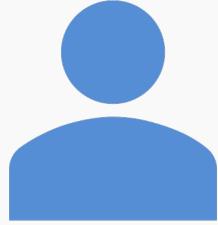


		It was spam, for real	It was not spam
		Four black envelope icons, each labeled "Spam".	One light blue envelope icon.
Filip: it is spam			
Filip: it is not spam		Two dark grey envelope icons, each labeled "Spam".	Three black envelope icons.

Generative Adversarial Networks (GANs)



Generative Adversarial Networks (GANs)



Generative Adversarial Networks (GANs)

Mary and Filip
**learned from what
went wrong from
their perspective**



	It was spam, for real	It was not spam
Filip: it is spam	Four icons of envelopes containing text, all labeled "Spam".	One icon of an envelope containing text.
Filip: it is not spam	One icon of an envelope containing text.	Four icons of envelopes containing text.

	It was spam, for real	It was not spam
Yes, it is spam	Four icons of envelopes containing text, all labeled "Spam".	One icon of an envelope containing text.
No, it is not spam	Two icons of envelopes containing text, labeled "Spam".	Two icons of envelopes containing text.

Generative Adversarial Networks (GANs)

Adversaries: Mary and Filip

Mary
Generator



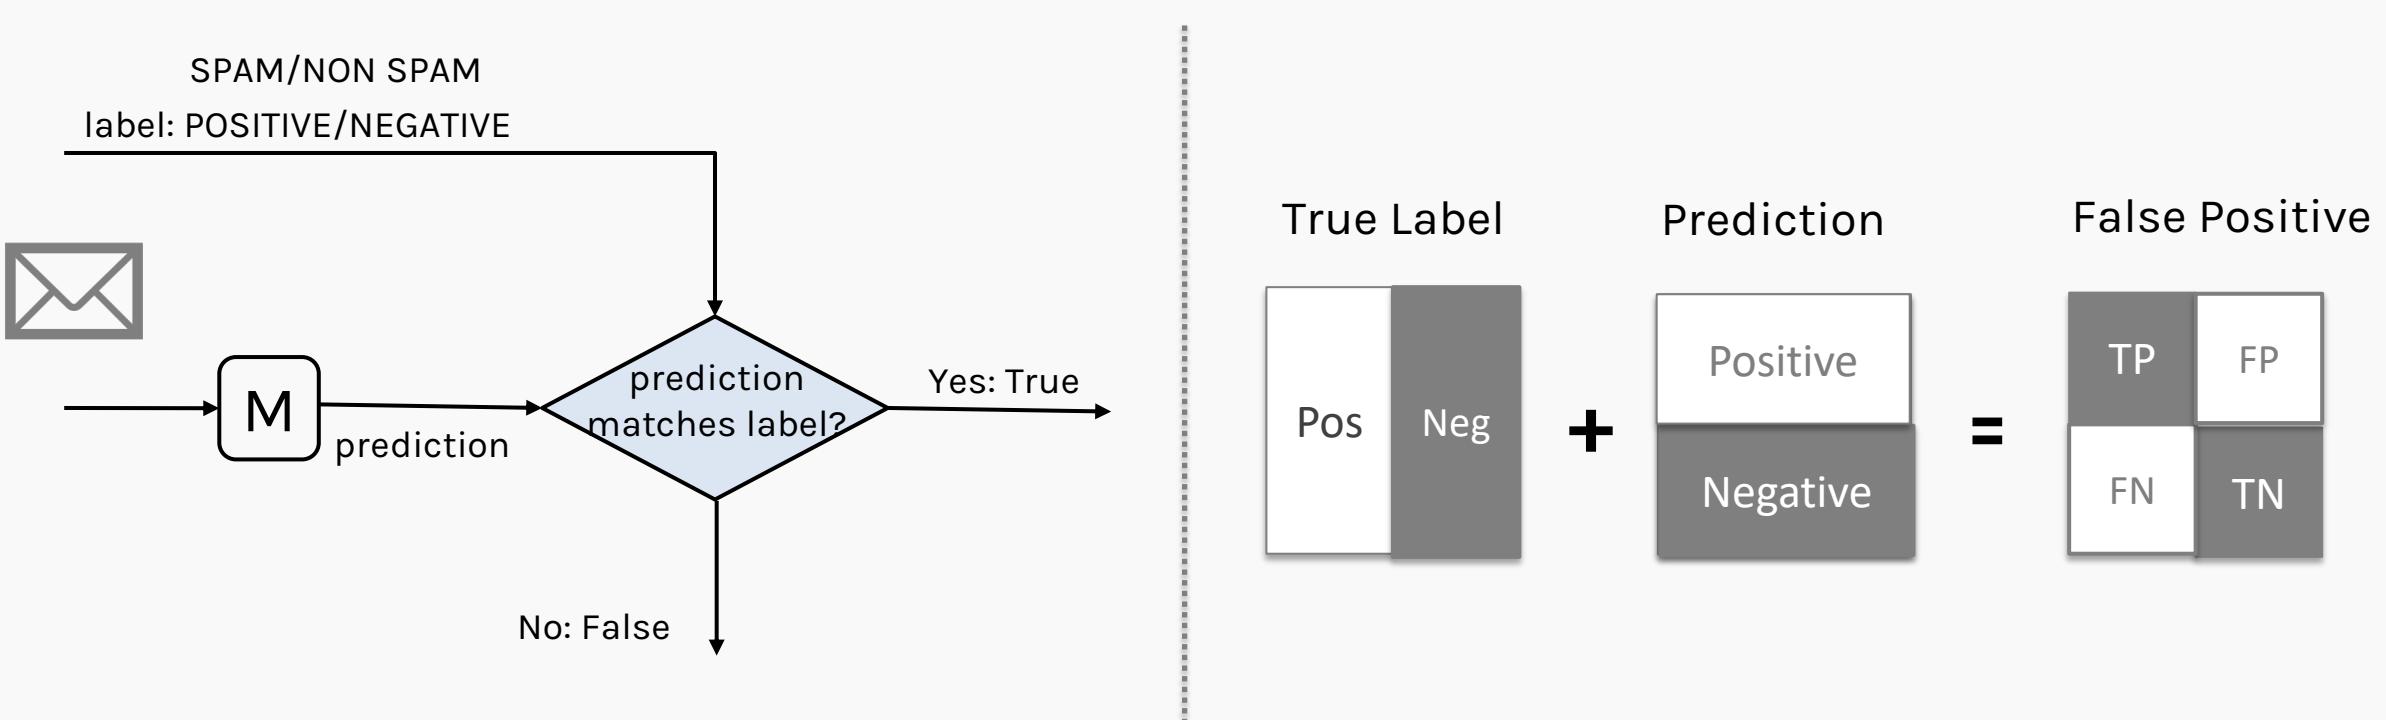
Filip
Discriminator



Generative Adversarial Networks (GANs)

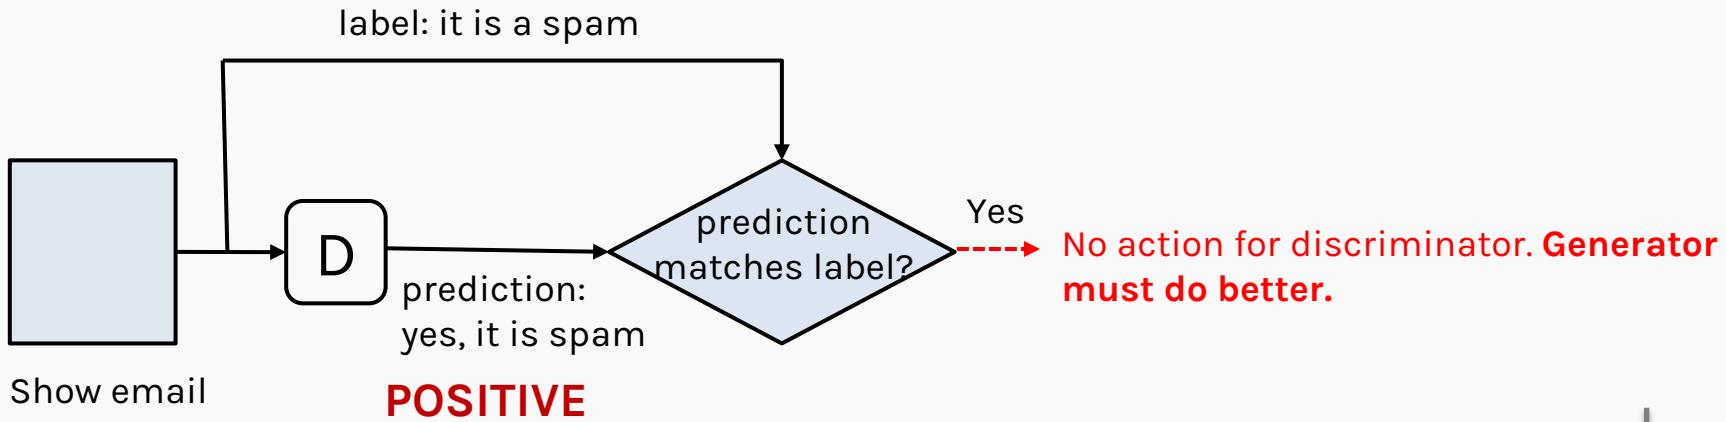
Understanding confusion matrix

TRUE/FALSE: If prediction and true label match / do not match
POSITIVE/NEGATIVE: Prediction class (SPAM = POSITIVE)



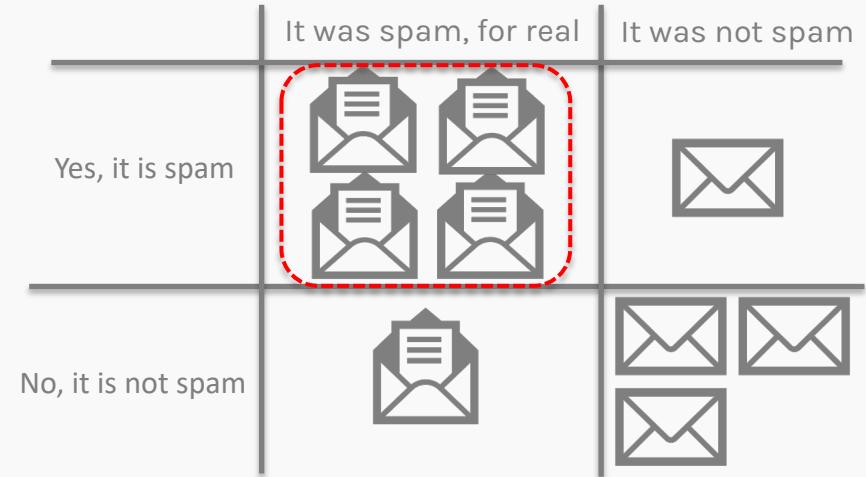
Generative Adversarial Networks (GAN)

TRUE/FALSE: If prediction and true label match / do not match
POSITIVE/NEGATIVE: Prediction class (SPAM = POSITIVE)



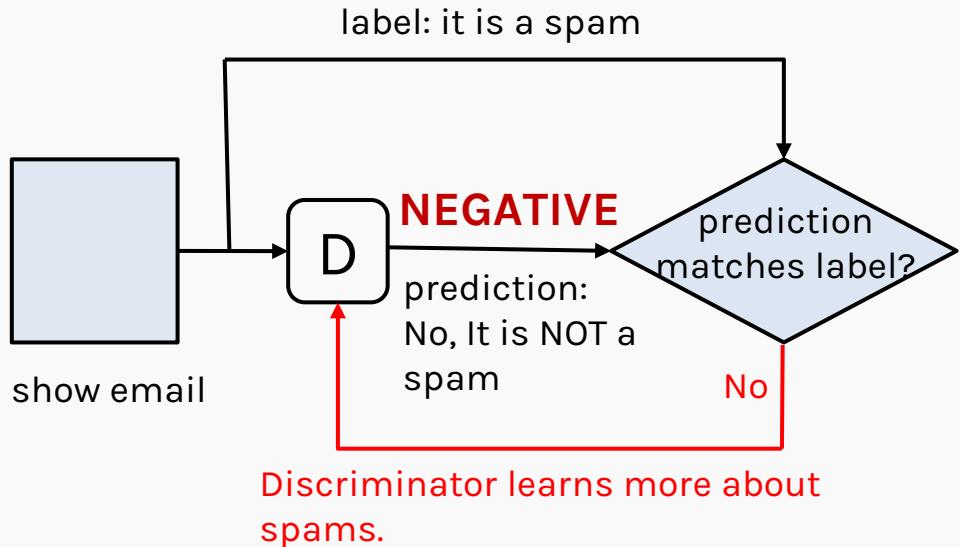
True positive (TP): the discriminator sees a spam and predicts correctly. No need for further actions for discriminator. Generator must do a better job.

TP	FP
FN	TN



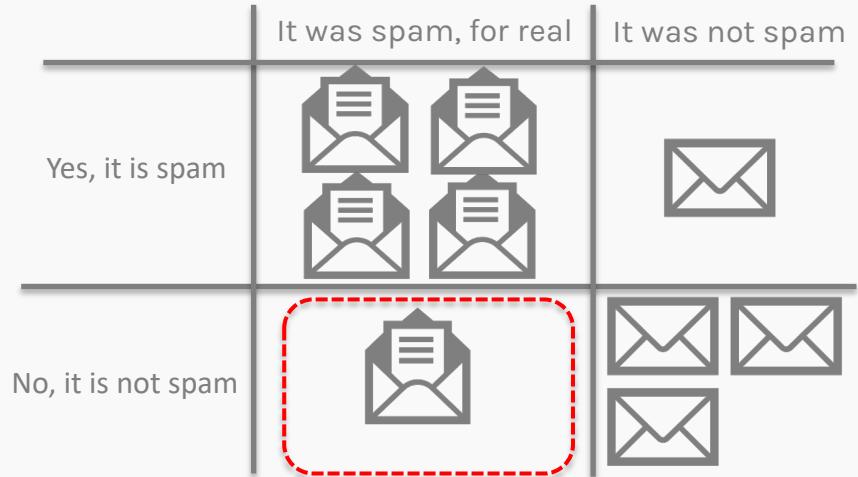
Generative Adversarial Networks (GANs)

TRUE/FALSE: If prediction and true label match / do not match
POSITIVE/NEGATIVE: Prediction class (SPAM = POSITIVE)

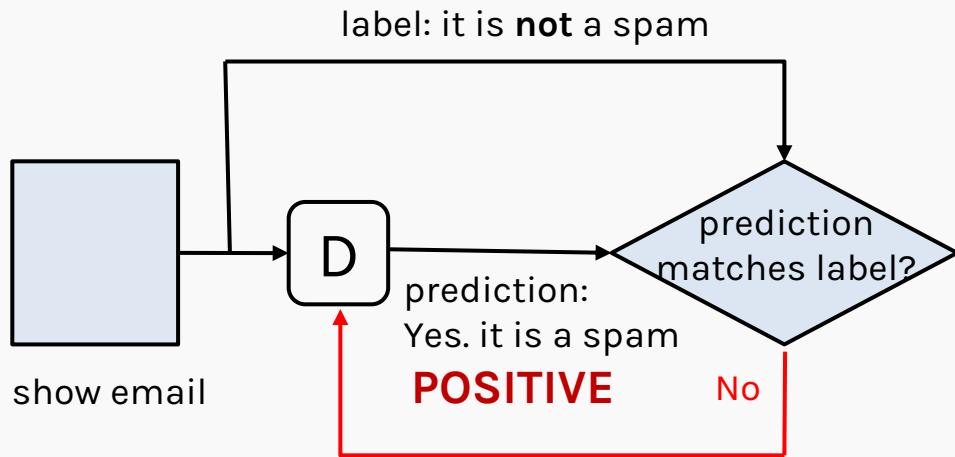


False Negative (FN): the discriminator sees an email and predicts it not a spam even though it is. The discriminator must learn more.

TP	FP
FN	TN



Generative Adversarial Networks (GANs)



Discriminator learns more about spams.

False Positive (FP): the discriminator sees an email and predicts it is a spam even though it is NOT. The discriminator must learn more.

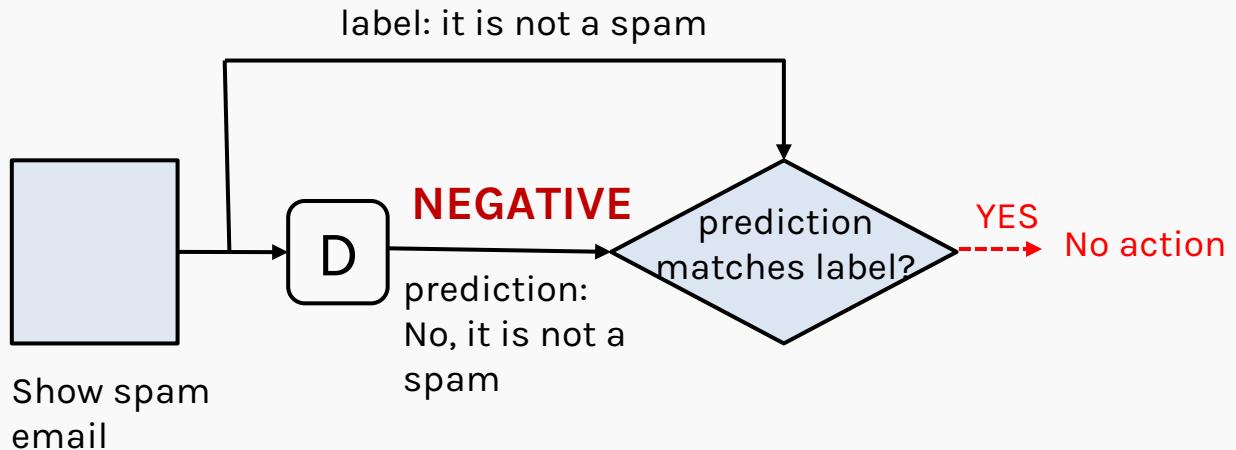
TP	FP
FN	TN

		It was spam, for real	It was not spam
Yes, it is spam	True Positives (TP)		
	False Positives (FP)		
No, it is not spam	False Negatives (FN)		
	True Negatives (TN)		

TRUE/FALSE: If prediction and true label match / do not match
POSITIVE/NEGATIVE: Prediction class (SPAM = POSITIVE)

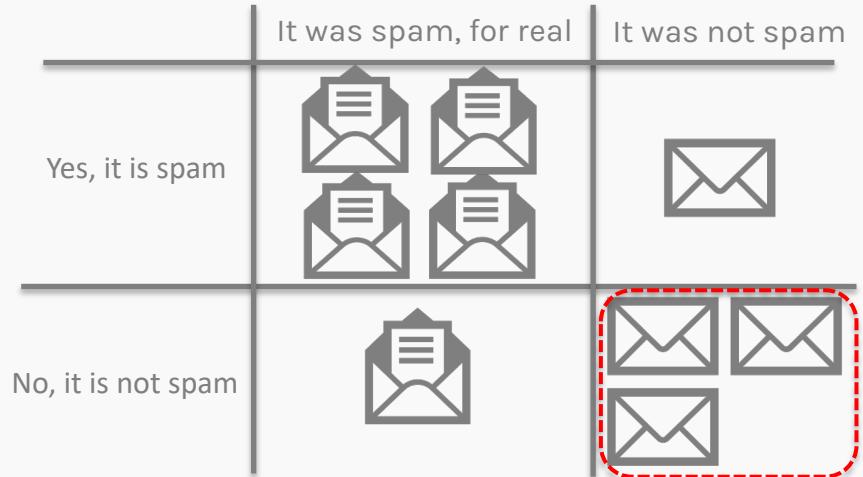
Generative Adversarial Networks (GANs)

TRUE/FALSE: If prediction and true label match / do not match
POSITIVE/NEGATIVE: Prediction class (SPAM = POSITIVE)



True negative (TN): No action required by Generator or Discriminator.

TP	FP
FN	TN



Generative Adversarial Networks (GANs)

Adversaries: Mary and Filip

Mary
Generator



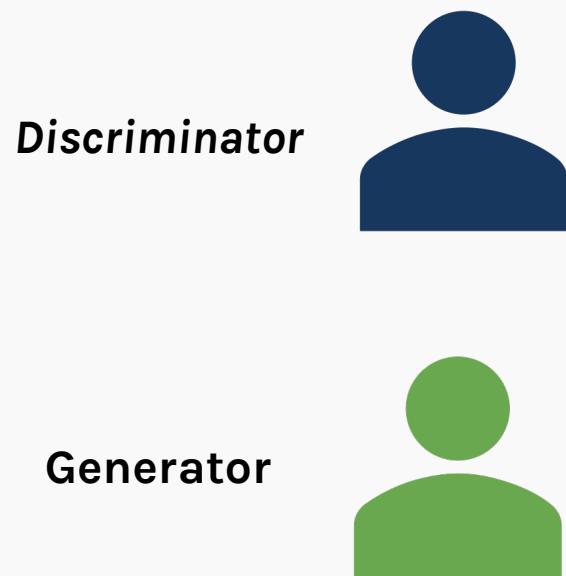
Filip
Discriminator



Generative Adversarial Networks (GANs)

Adversaries: Mary and Filip

Become: Two player game between a **generator G** and a **discriminator D**.



Why is it a “game” ?



Outline

- Generative Modeling Motivation
- High Level Formalism
- **Architecture**
- Mathematics
- Training GANS
- Conditional GANS

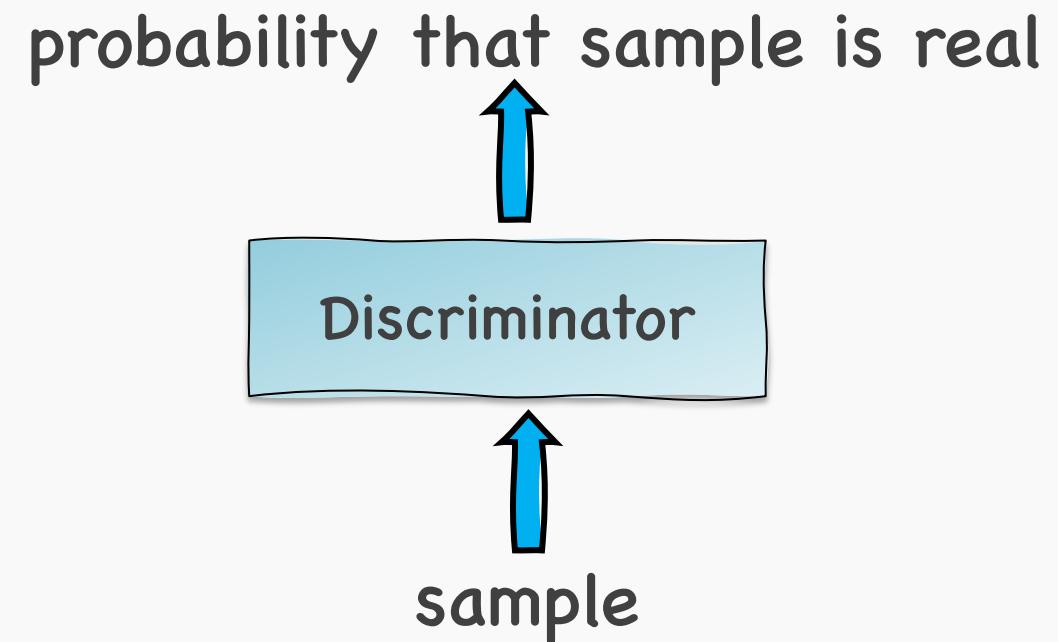


The Discriminator



The **discriminator** is very simple. It takes a sample as input, and its output is a single value that reports the network's probability that the input is from the training set rather than being a fake.

There are not many restrictions on what the discriminator is.



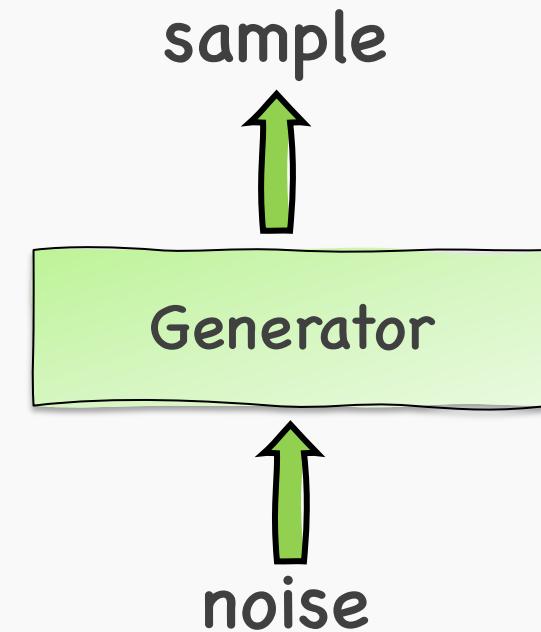
The Generator



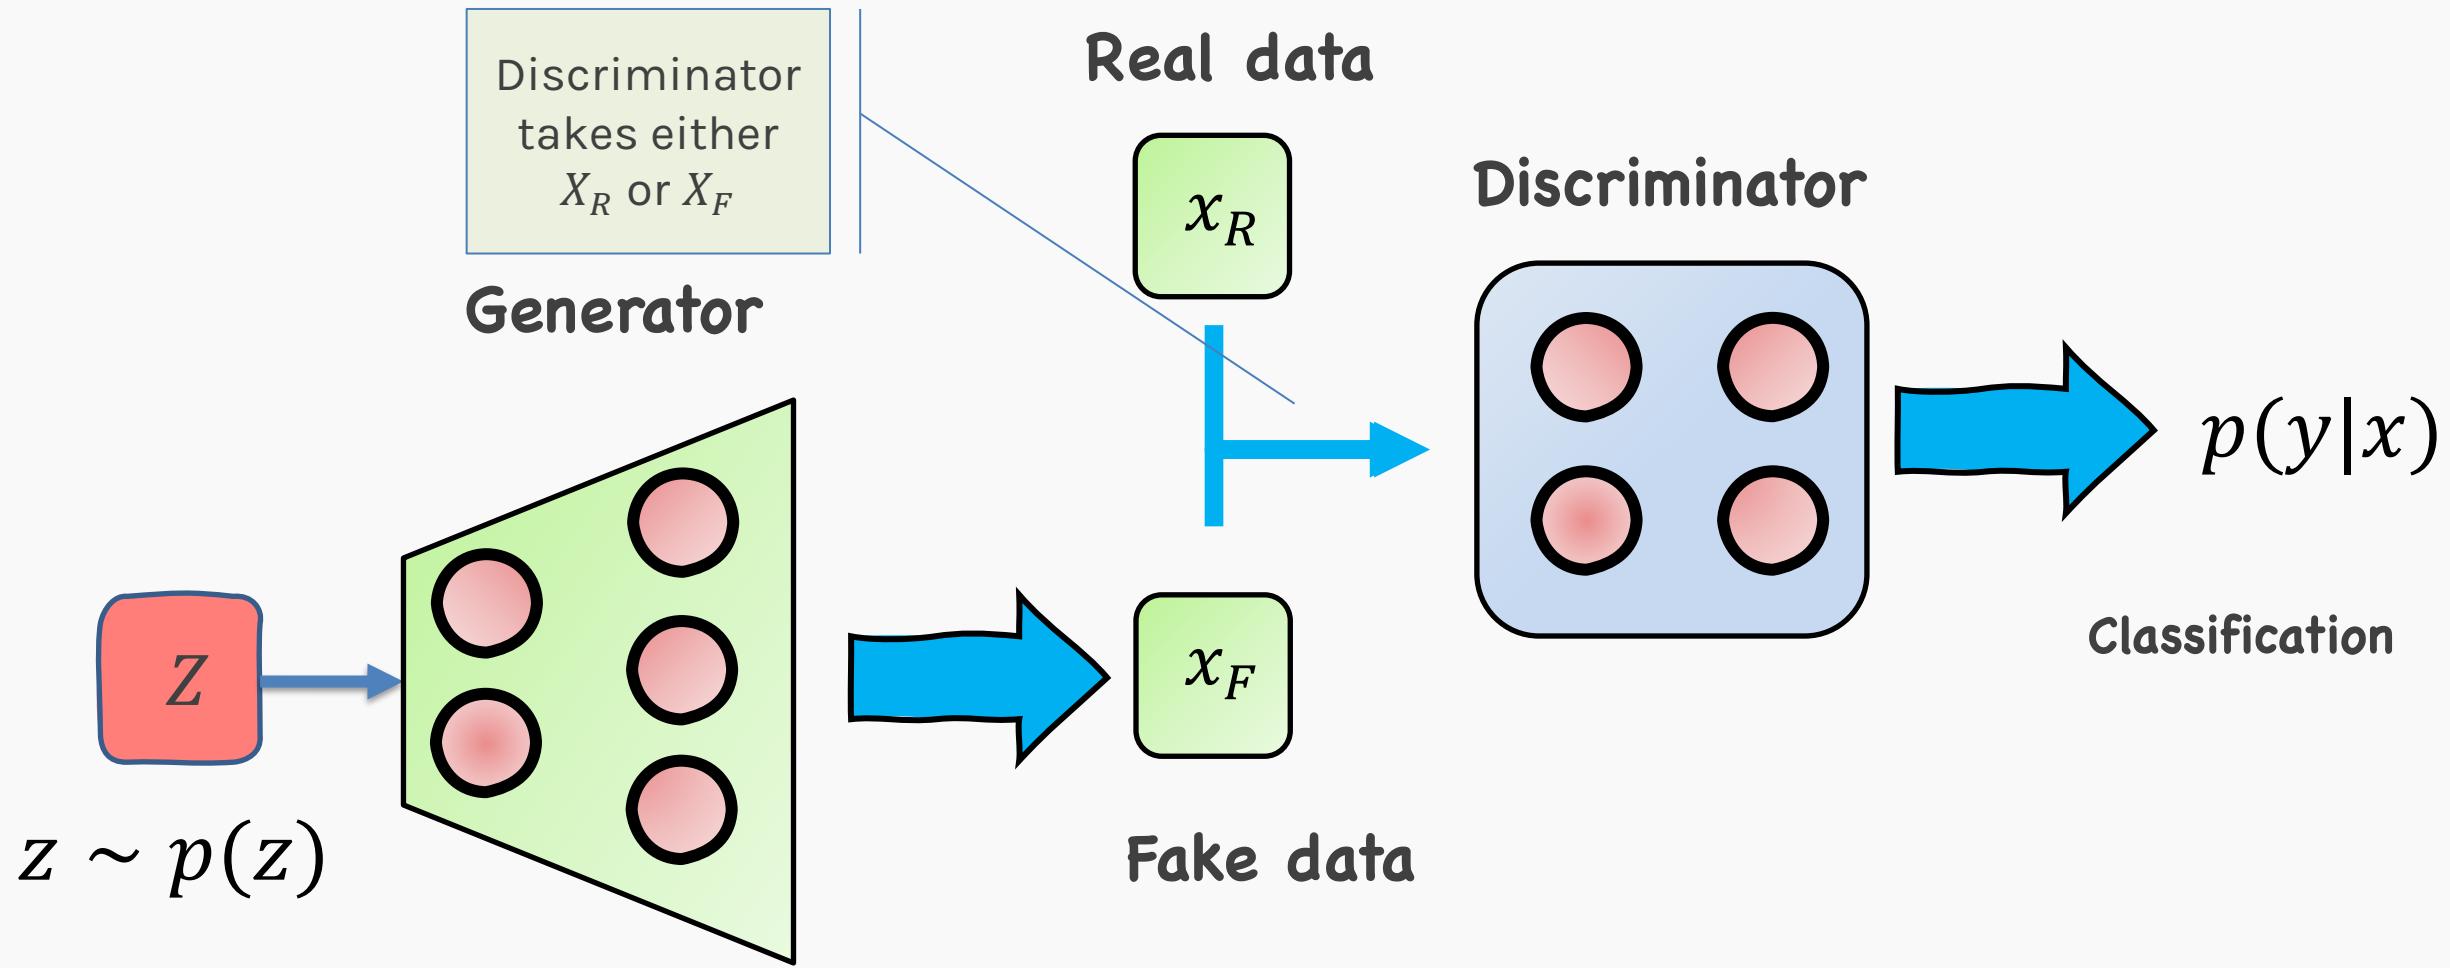
The **generator** takes as input a bunch of random numbers and generates a sample.

If we build our generator to be deterministic, then the same input will always produce the same output.

We want to generate data from a **distribution**. In that sense, we can think of the input values as latent variables.



GANS Architecture



Outline

- Generative Modeling Motivation
- High Level Formalism
- Architecture
- **Mathematics**
- Training GANS
- Conditional GANS



Training: Loss Function

In a binary classification problem, the loss function is given by:

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N y_i \log(p(y_i)) + (1 - y_i) \log(1 - p(y_i))$$

Where y is the label for *real*=1 or *fake*=0. The input to the **Discriminator** can be the real data or the fake data generated by the **Generator**. Splitting the loss function, we have:

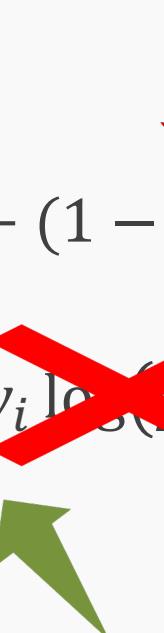
$$\mathcal{L} = -\frac{1}{N_{Real}} \sum_{i=1}^{N_{Real}} y_i \log(p(y_i)) + (1 - y_i) \log(p(1 - y_i)) - \frac{1}{N_{Fake}} \sum_{i=1}^{N_{Fake}} y_i \log(p(y_i)) + (1 - y_i) \log(1 - p(y_i))$$

Learning

$$\mathcal{L} = -\frac{1}{N_{Real}} \sum_{i=1}^{N_{Real}} \cancel{y_i \log(p(y_i)) + (1 - y_i) \log(1 - p(y_i))}$$
$$-\frac{1}{N_{Fake}} \sum_{i=1}^{N_{Fake}} \cancel{y_i \log(p(y_i)) + (1 - y_i) \log(1 - p(y_i))}$$

Real: $y_i = 1$

Fake: $y_i = 0$



Learning

$$\mathcal{L} = -\frac{1}{N_{Real}} \sum_{i=1}^{N_{Real}} \log(p(y_i)) - \frac{1}{N_{Fake}} \sum_{i=1}^{N_{Fake}} \log(1 - p(y_i))$$

Learning

$$\mathcal{L} = -\frac{1}{N_{Real}} \sum_{i=1}^{N_{Real}} \log(p(y_i)) - \frac{1}{N_{Fake}} \sum_{i=1}^{N_{Fake}} \log(1 - p(y_i))$$

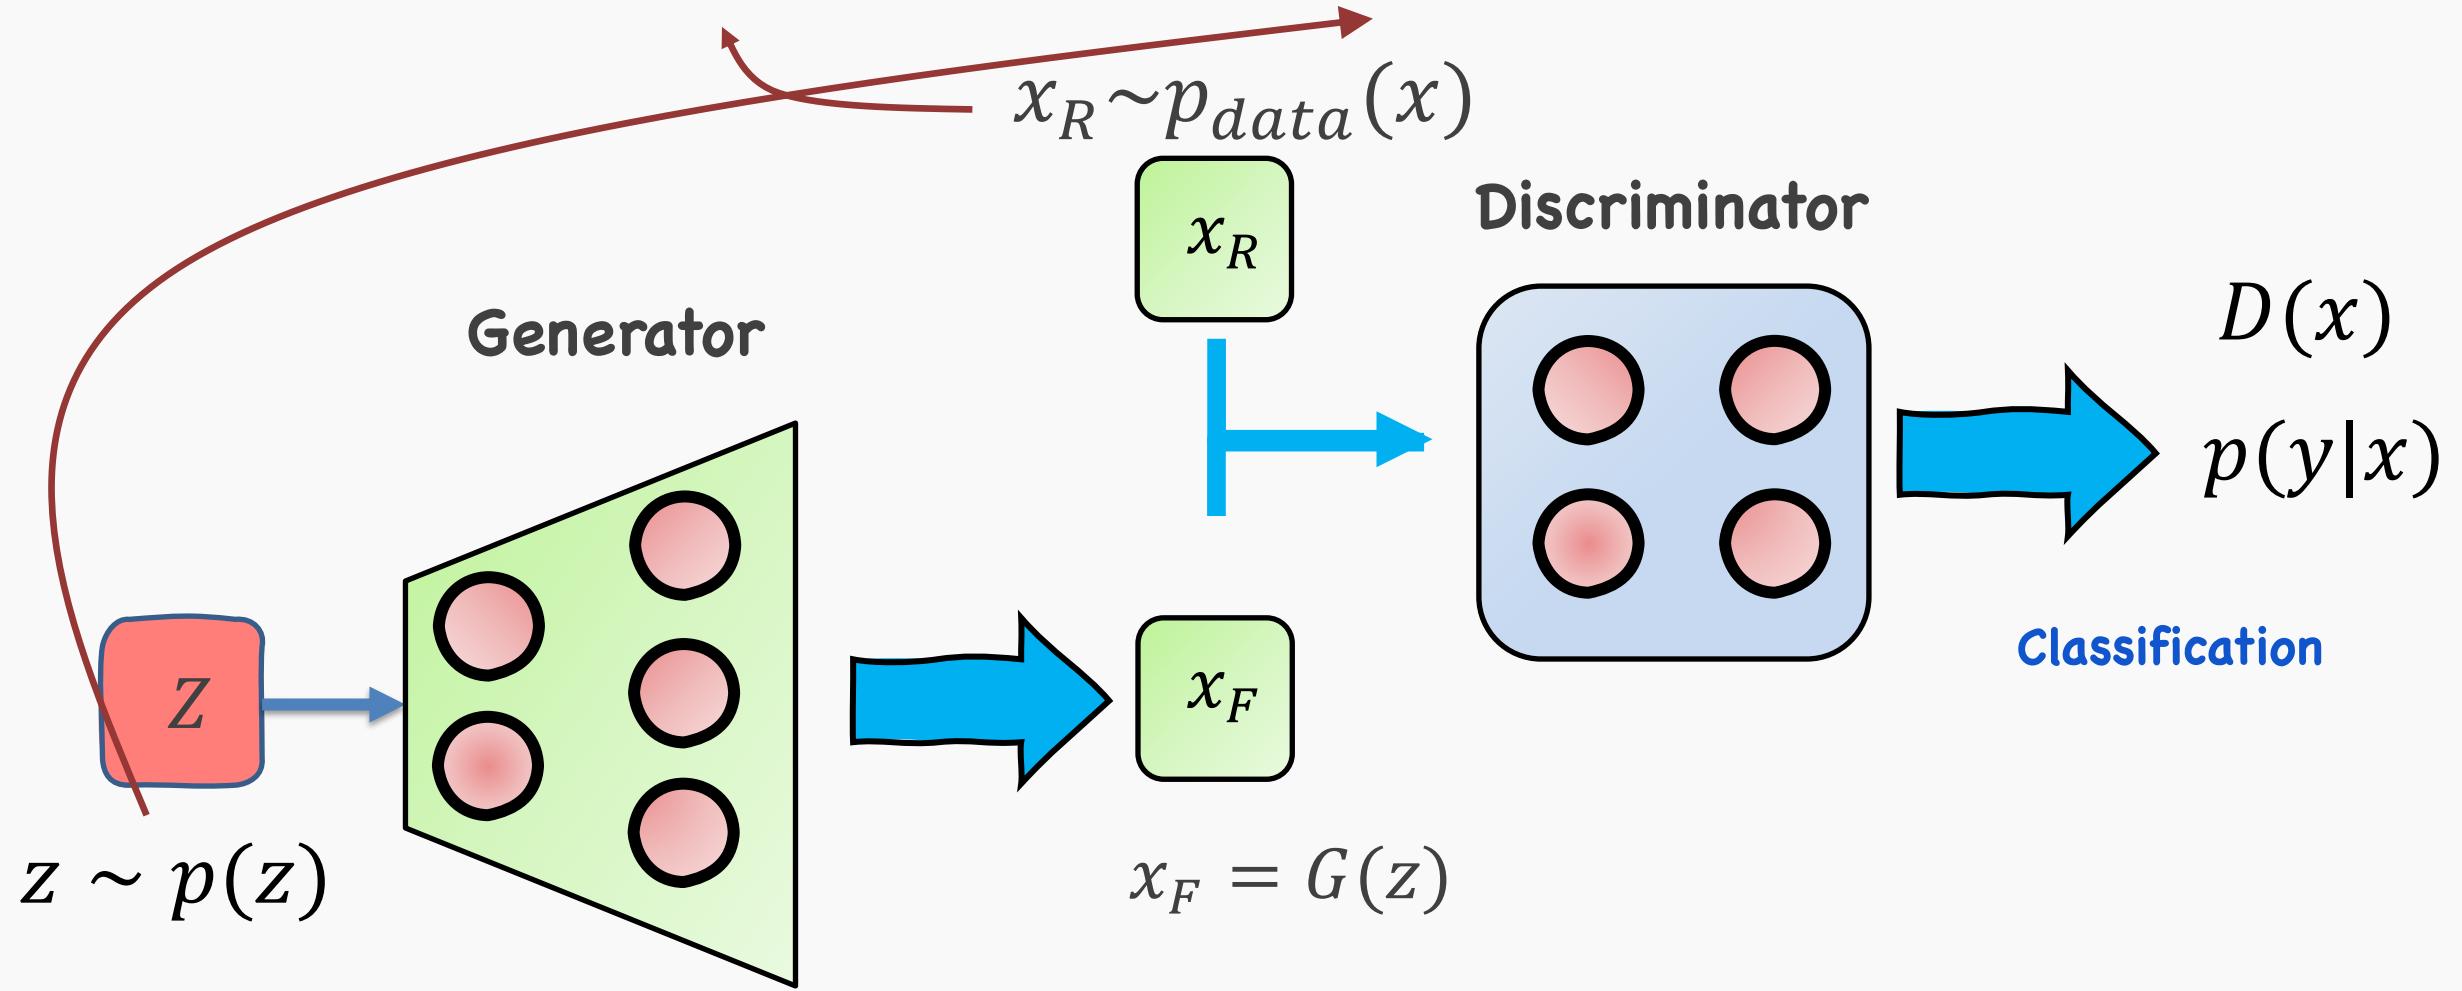
Rewriting in terms of discriminator **D** and generator **G** outputs:

$$\mathcal{L} = -\frac{1}{N_{Real}} \sum_{i=1}^{N_{Real}} \log(D(x_i^R)) - \frac{1}{N_{Fake}} \sum_{i=1}^{N_{Fake}} \log(1 - D(x_i^F))$$

And noting that $x_i^F = G(z_i)$

$$\mathcal{L} = -\frac{1}{N_{Real}} \sum_{i=1}^{N_{Real}} \log(D(x_i^R)) - \frac{1}{N_{Fake}} \sum_{i=1}^{N_{Fake}} \log(1 - D(G(z_i)))$$

$$\mathcal{L} = -\frac{1}{N_{Real}} \sum_{i=1}^{N_{Real}} \log(D(x_i^R)) - \frac{1}{N_{Fake}} \sum_{i=1}^{N_{Fake}} \log(1 - D(G(z_i)))$$



$$\mathcal{L} = -\mathbb{E}_{x \sim p_{data}(x)}[\log(D(x))] - \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$$

Learning

$$\mathcal{L} = -\mathbb{E}_{x \sim p_{data}(x)}[\log(D(x))] - \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$$

The adversarial training can be described as though the **Generator G** and **Discriminator D** play the following two-player min-max game with the following value function V (G, D).

The **Discriminator's job** is to minimize the loss or maximize the -ve loss.

$$\max_D V(G, D) = \mathbb{E}_{x \sim p_{data}(x)}[\log(D(x))] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$$

The **Generator's job** is to maximize the loss or minimize the -ve loss.

$$\min_G \max_D V(G, D) = \mathbb{E}_{x \sim p_{data}(x)}[\log(D(x))] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$$

Minimax value function

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

Discriminator's prediction on real data

Discriminator's prediction on fake data

Sample real data

Sample random noise

Generator's output:
fake data

**Minimax value
function**

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

**Discriminator:
Maximize**

**Generator:
Minimize**

**Minimax value
function**

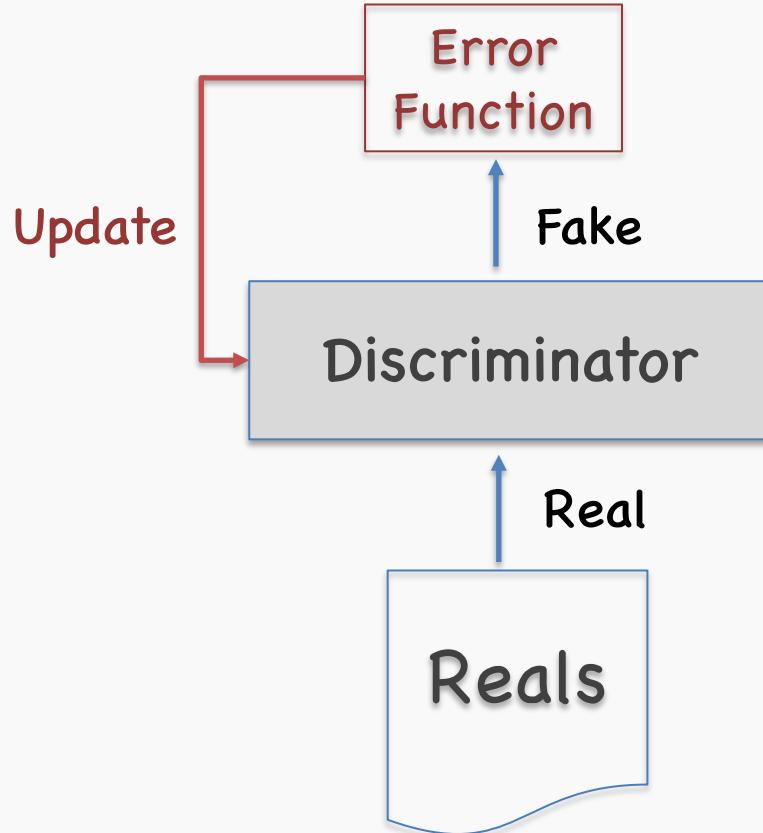
$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

**Discriminator:
Maximize**

**Generator:
Maximize**

$$\mathbb{E}_{z \sim p_z(z)} [\log(D(G(z)))]$$

Training the GAN



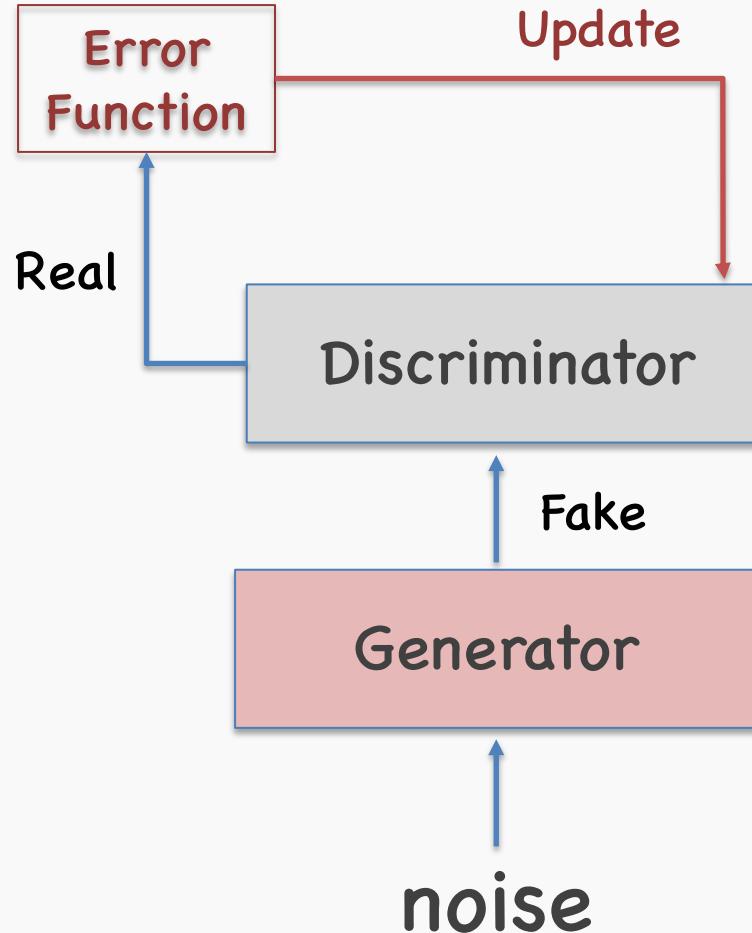
$$\max_D \mathbb{E}_{x \sim p_{data}(x)} [\log(D(x))]$$

False negative (I: Real/D: Fake):

In this case we feed reals to the discriminator. The Generator is not involved in this step at all.

The error function here only involves the Discriminator and if it makes a mistake the error drives a backpropagation step through the discriminator, updating its weights, so that it will get better at recognizing reals.

Training the GAN

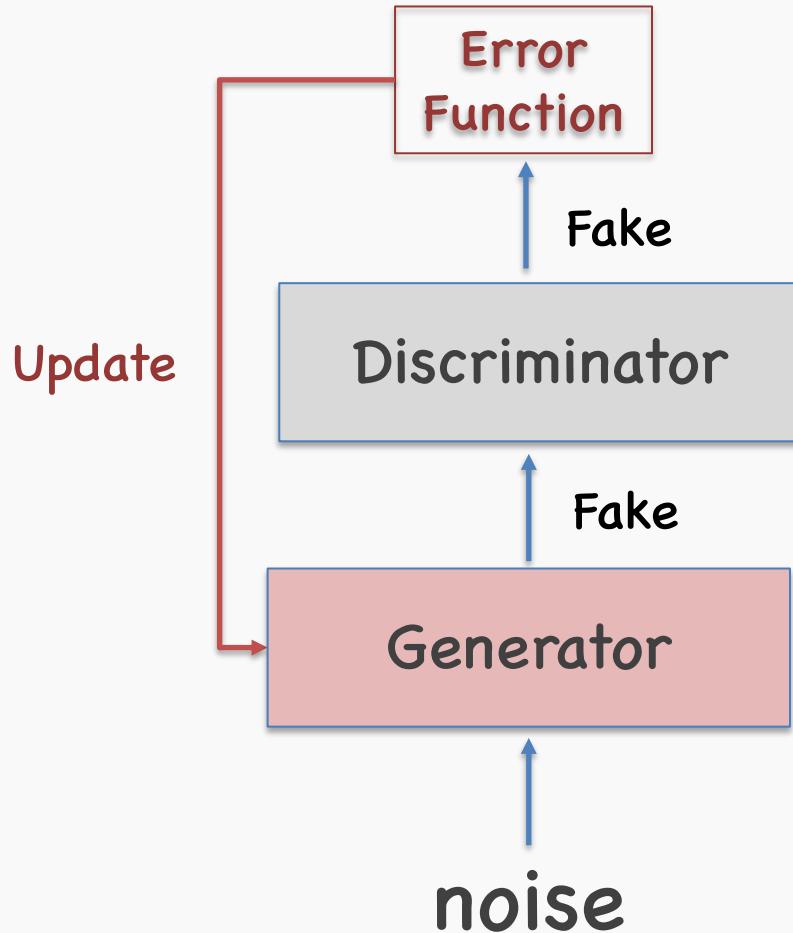


$$\max_D \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

False positives (I:Fake/D:Real):

Here we generate a fake and punish the discriminator if it classifies it as real.

Training the GAN



$$\max_G \mathbb{E}_{z \sim p_z(z)} [\log(D(G(z)))]$$

True negative (I: Fake/D: Fake):

- We start with random numbers going into the generator.
- The generator's output is a fake.
- The objective function gets a large -ve value if this fake is correctly identified as fake. Meaning that the generator got caught.
- Backprop, goes through the discriminator (which is frozen) to the generator updating the generator's weight, so it can better learn how to fool the discriminator.

Learning

The process – known as **Learning Round** - accomplishes three jobs:

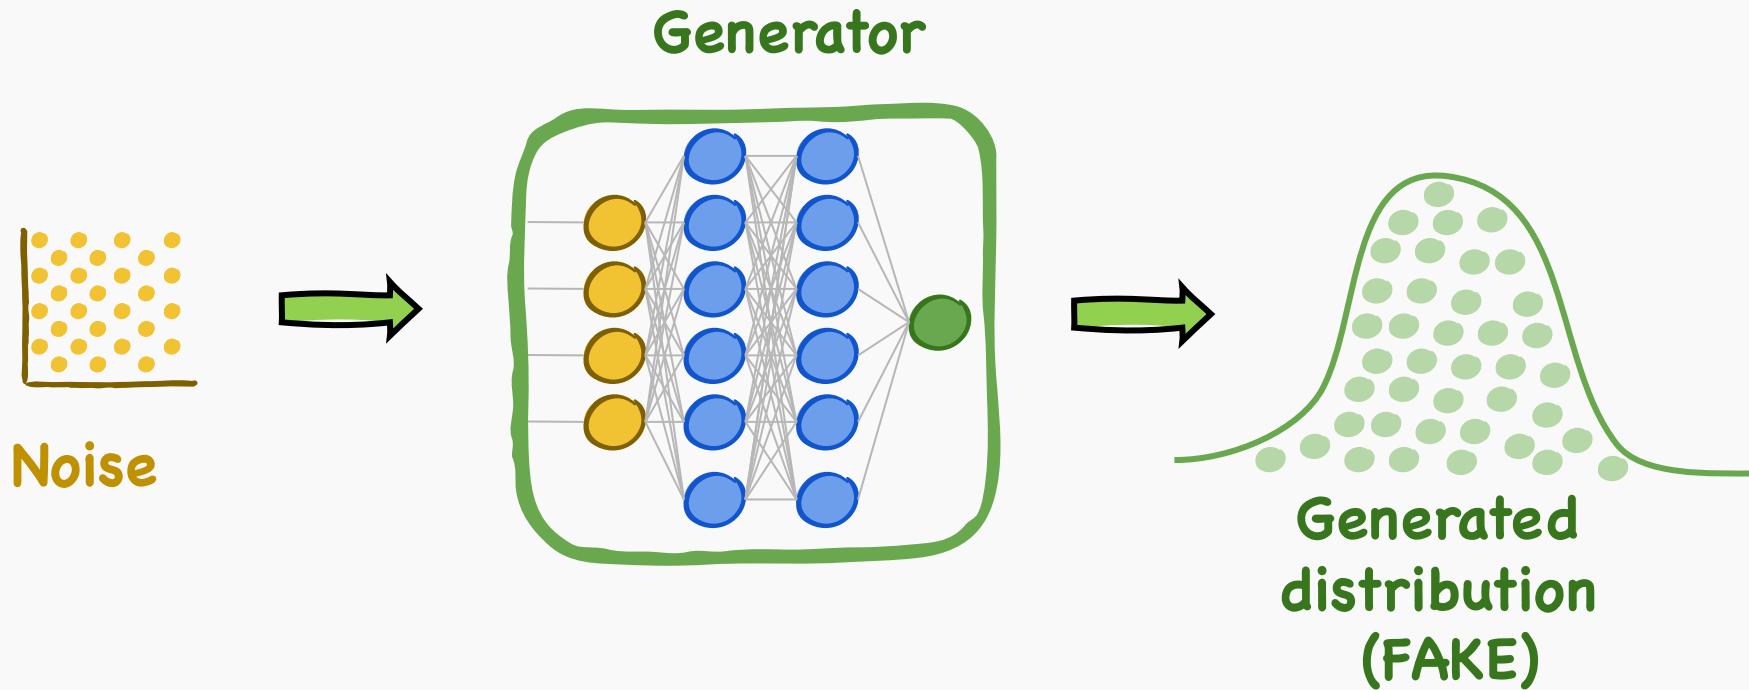
1. The discriminator learns to identify features that characterize a real sample
2. The discriminator learns to identify features that reveal a fake sample
3. The generator learns how to avoid including the features that the discriminator has learned to spot

Outline

- Generative Modeling Motivation
- High Level Formalism
- Architecture
- Mathematics
- **Training GANS**
- Conditional GANS

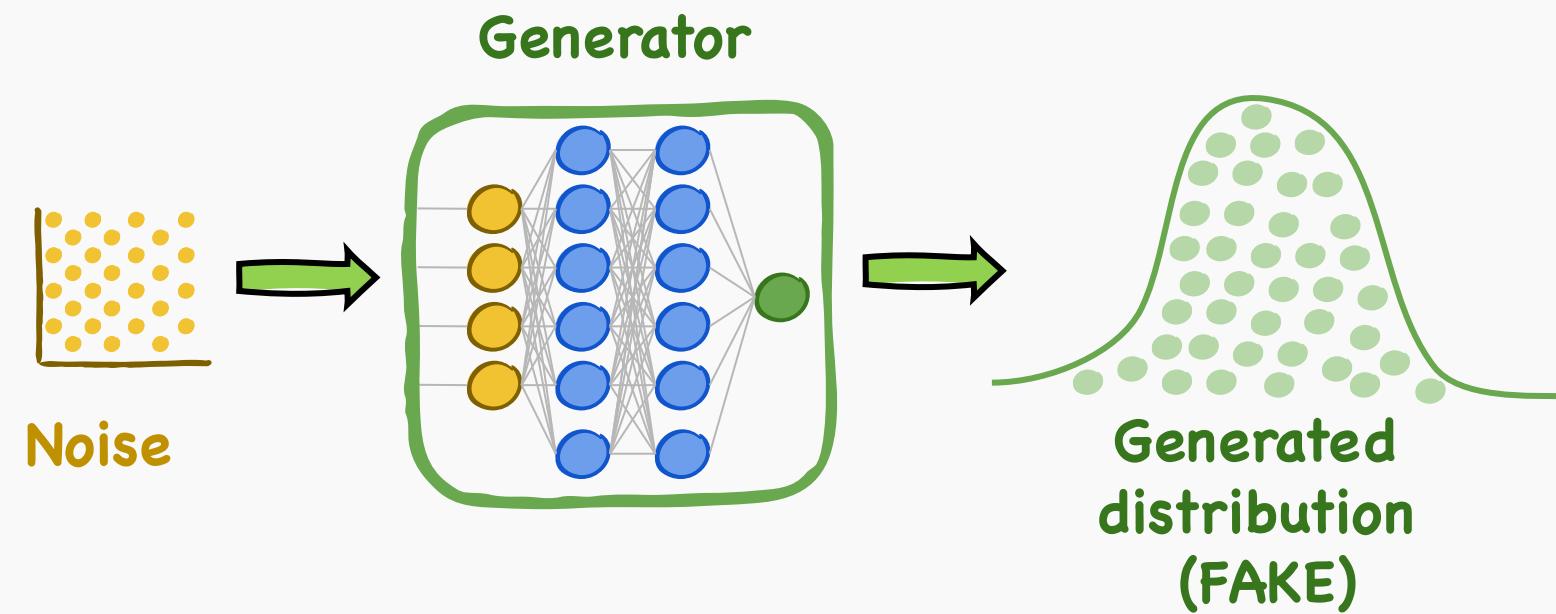


The Generator



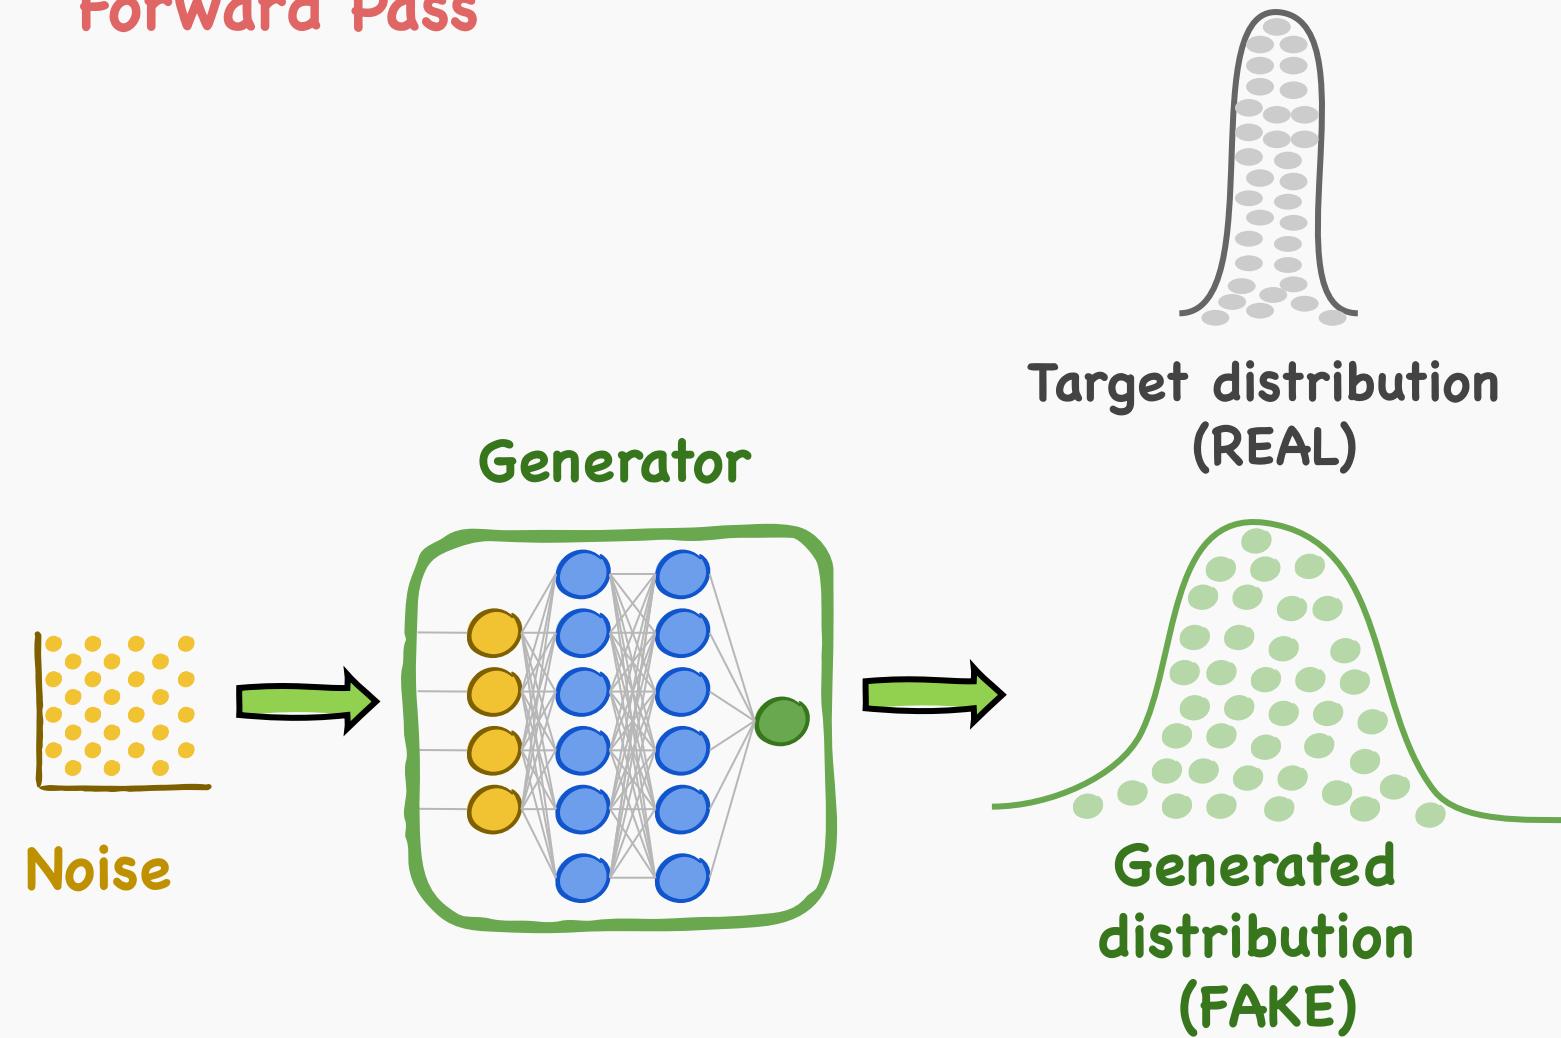
Training GANs

Forward Pass



Training GANs

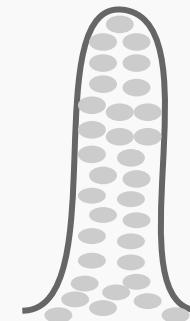
Forward Pass



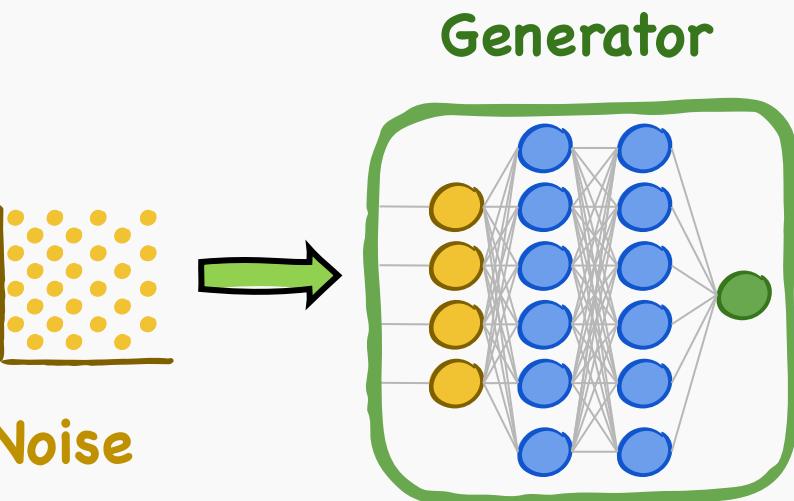
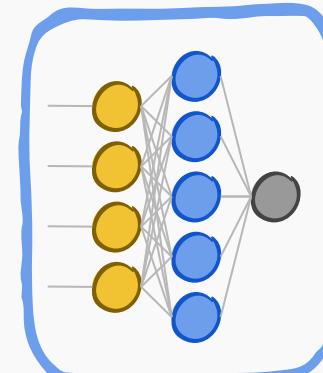
Training GANs

Forward Pass

Target distribution
(REAL)



Discriminator

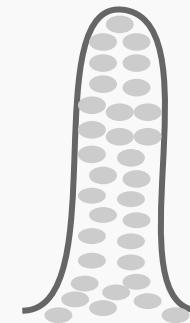


Generated distribution
(FAKE)

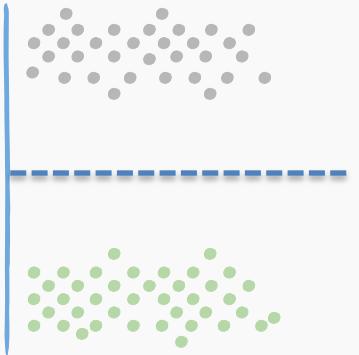
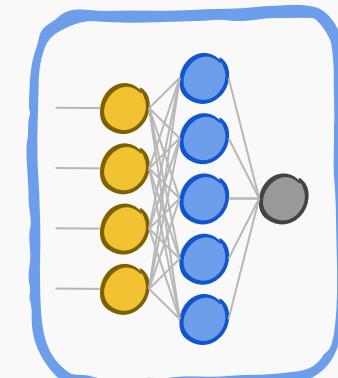
Training GANs

Forward Pass

Target distribution
(REAL)

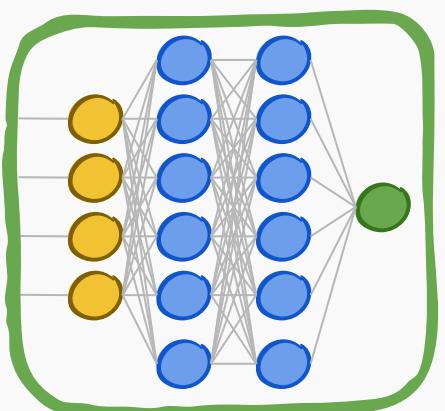


Discriminator



Real vs. Fake
Classification

Generator



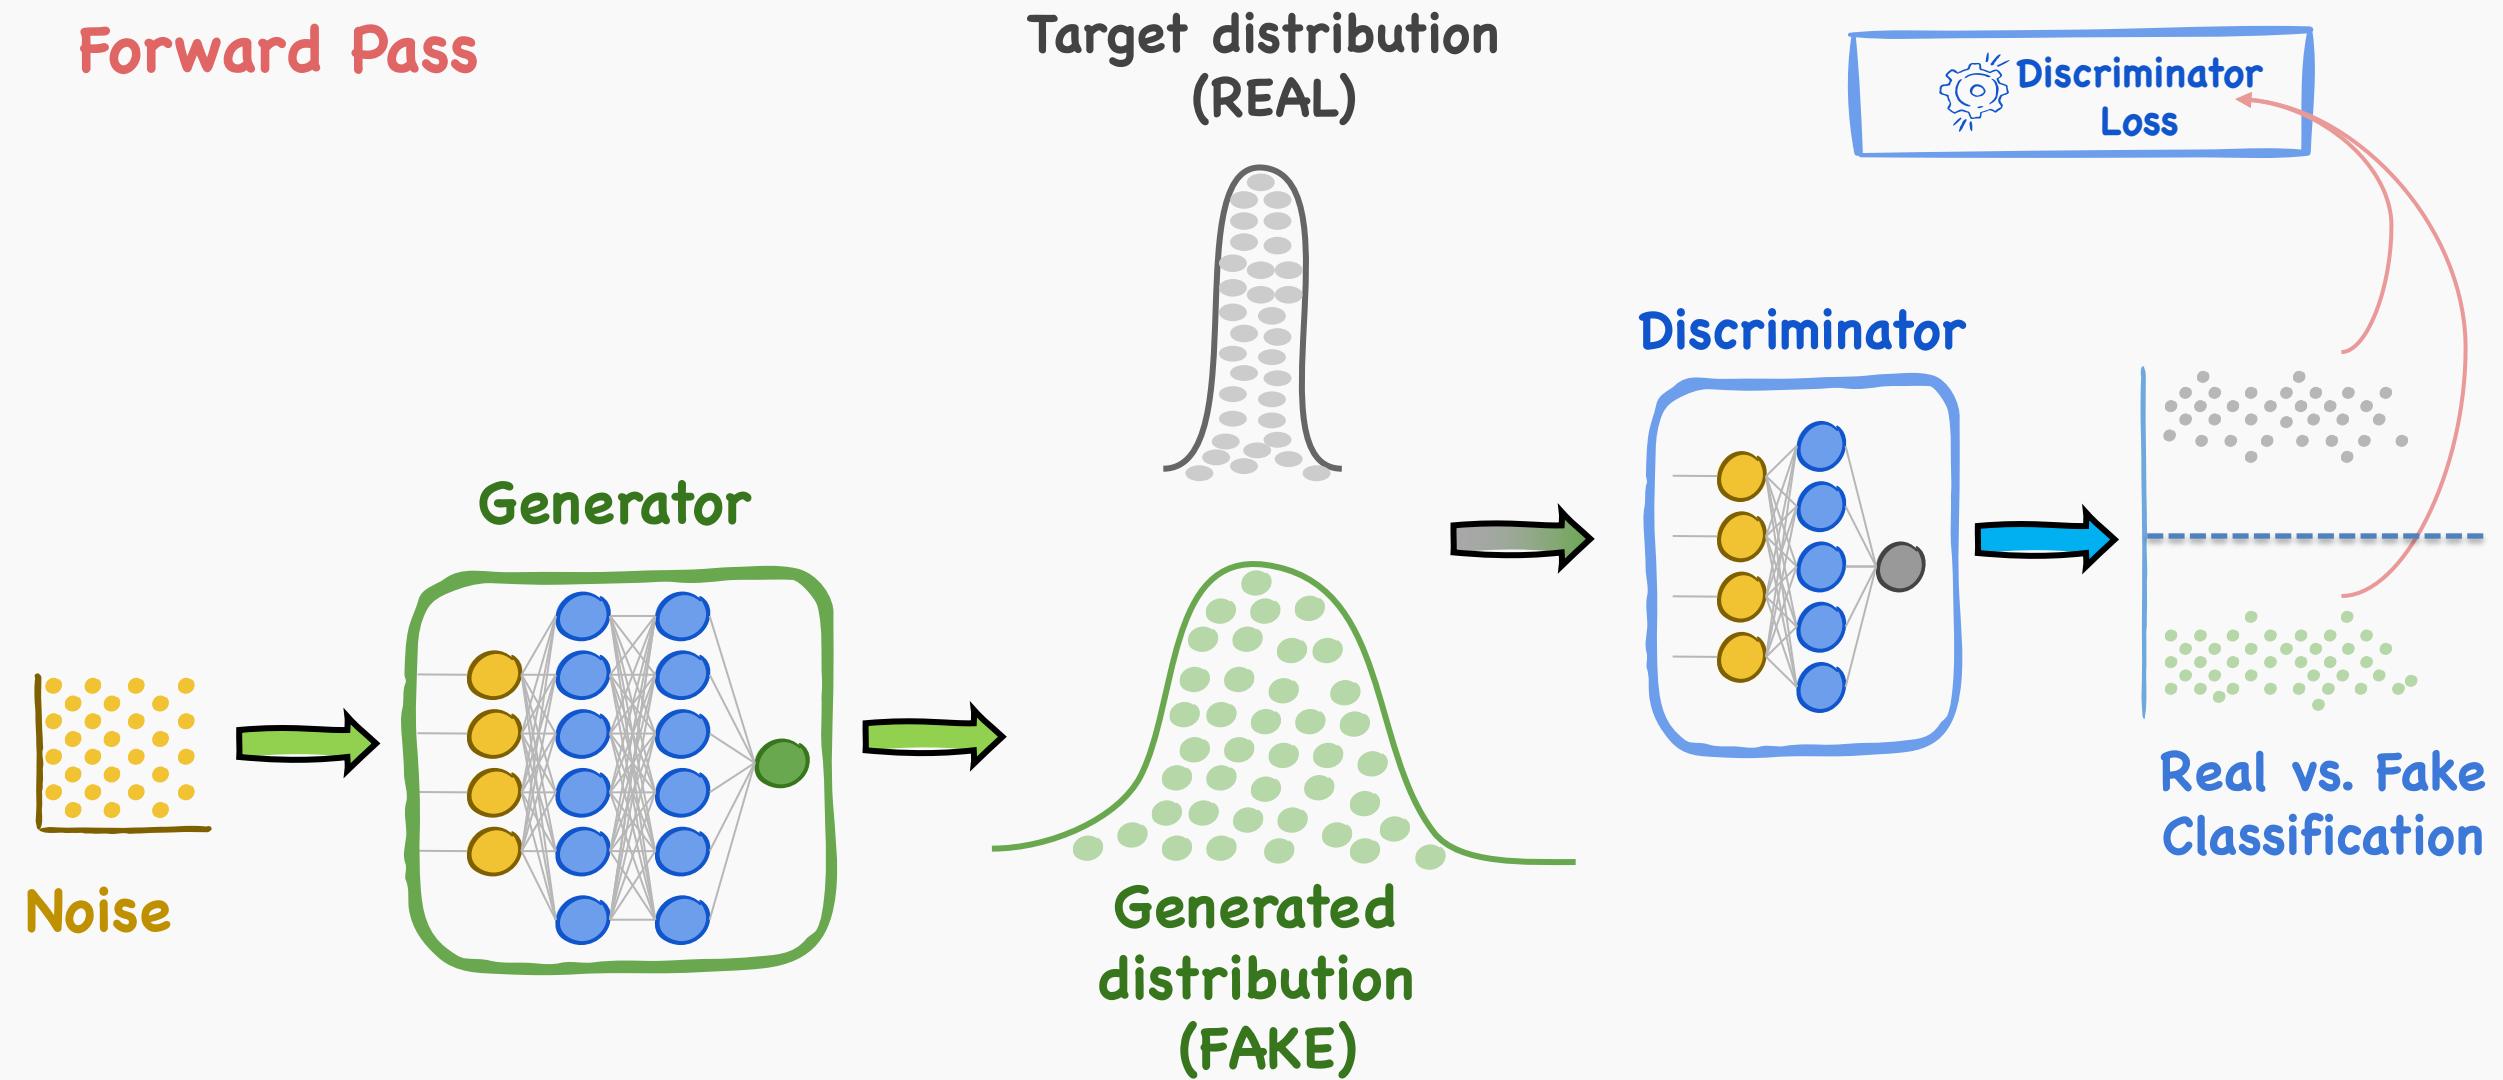
Generated
distribution
(FAKE)

Noise



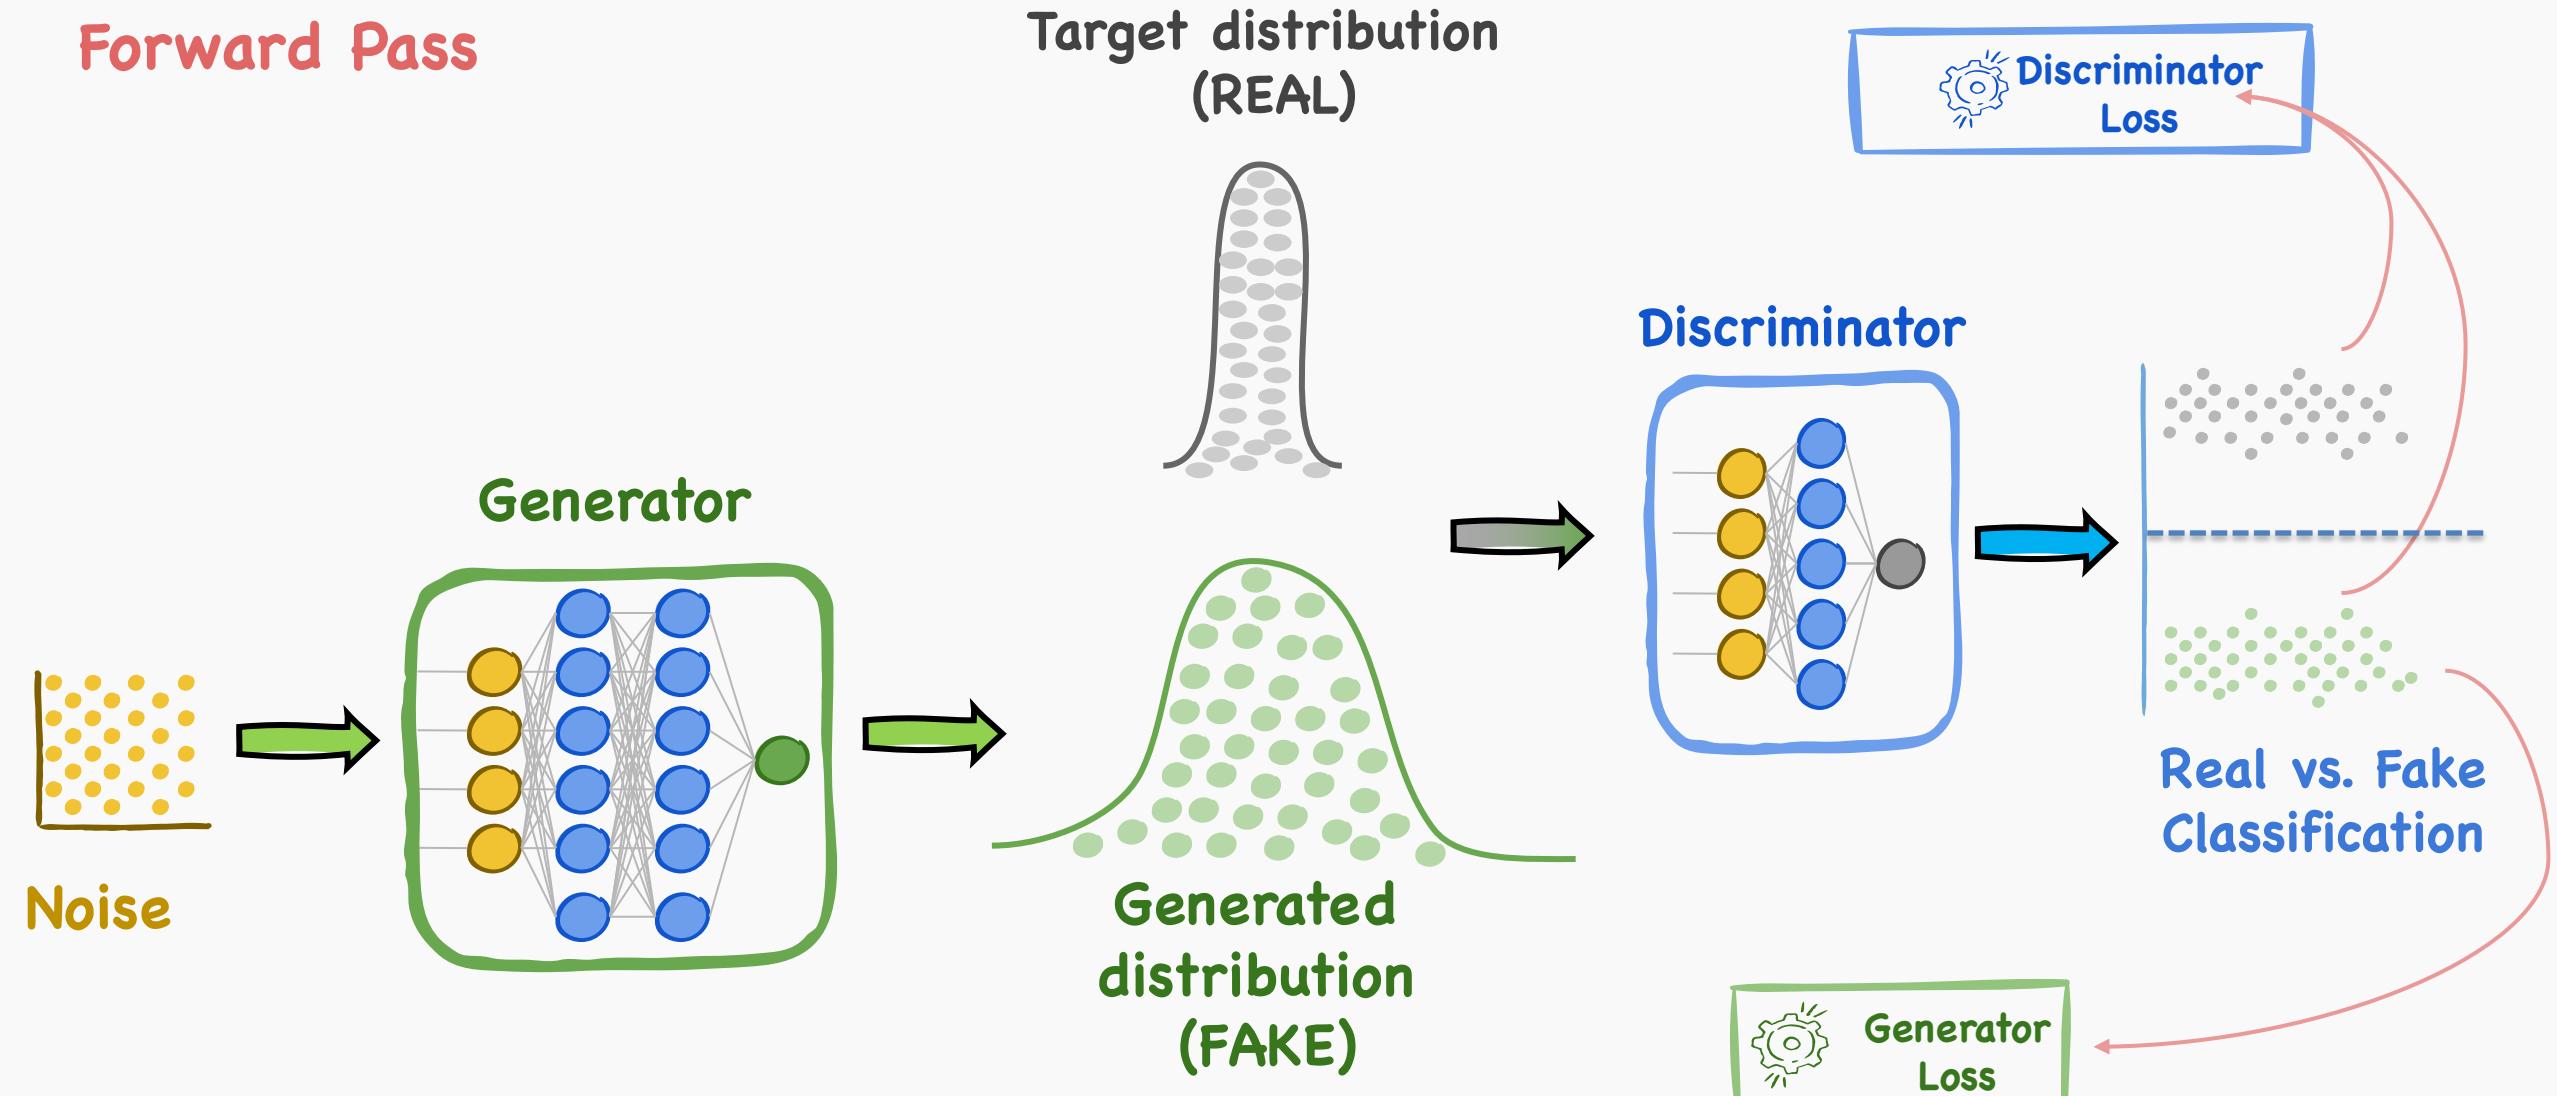
Training GANs

Forward Pass



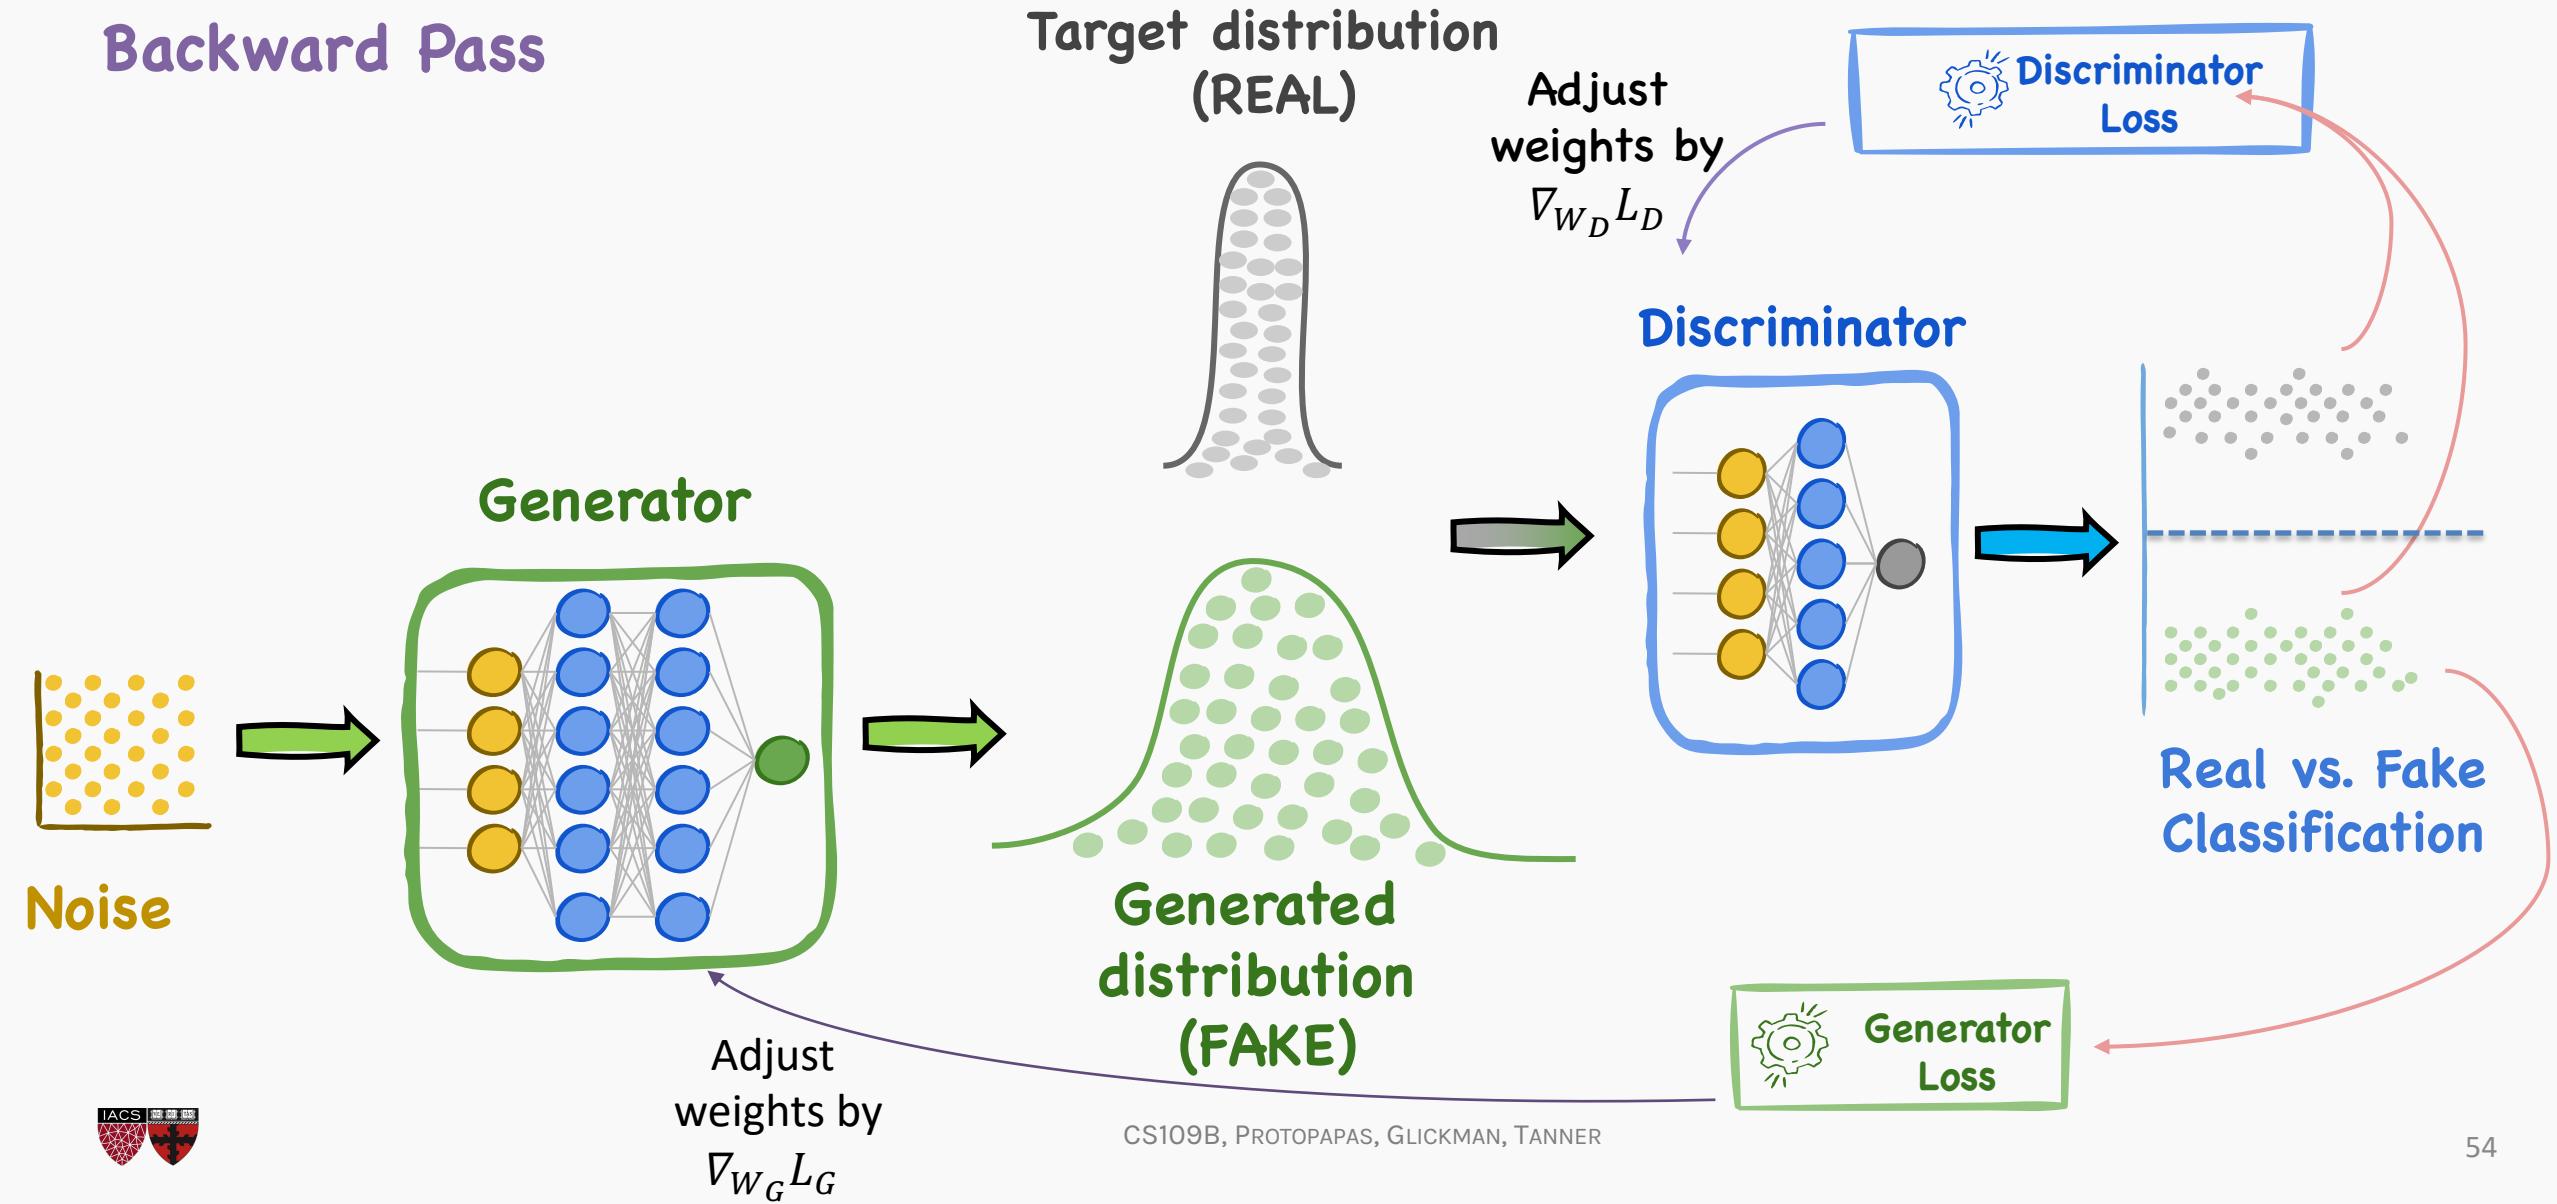
Training GANs

Forward Pass



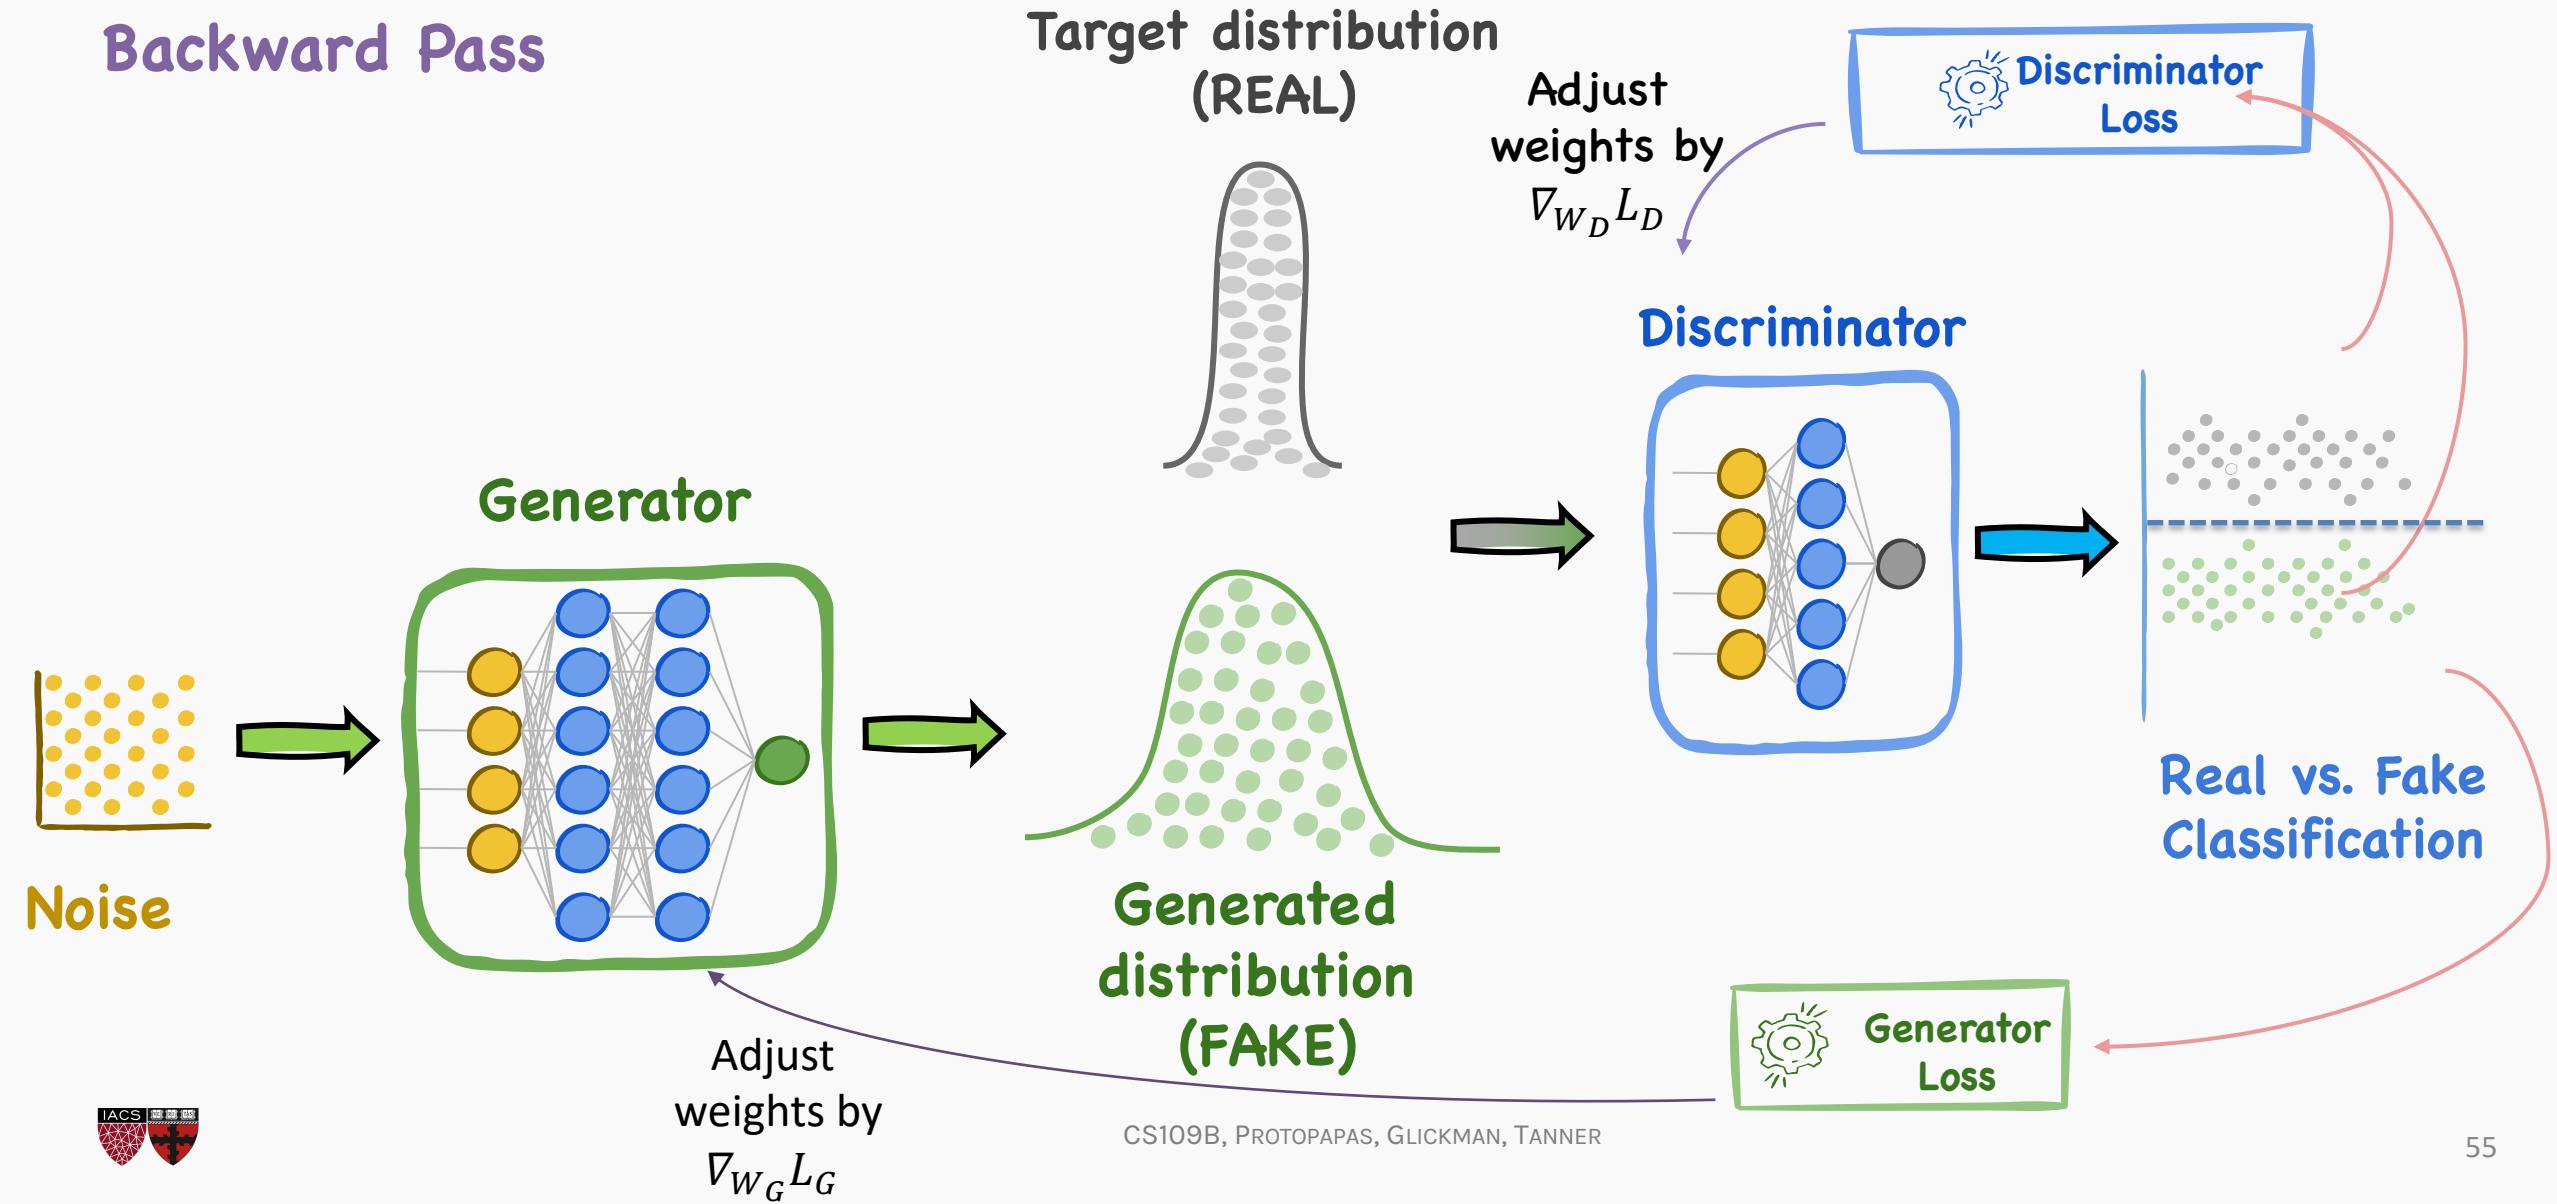
Training GANs

Backward Pass



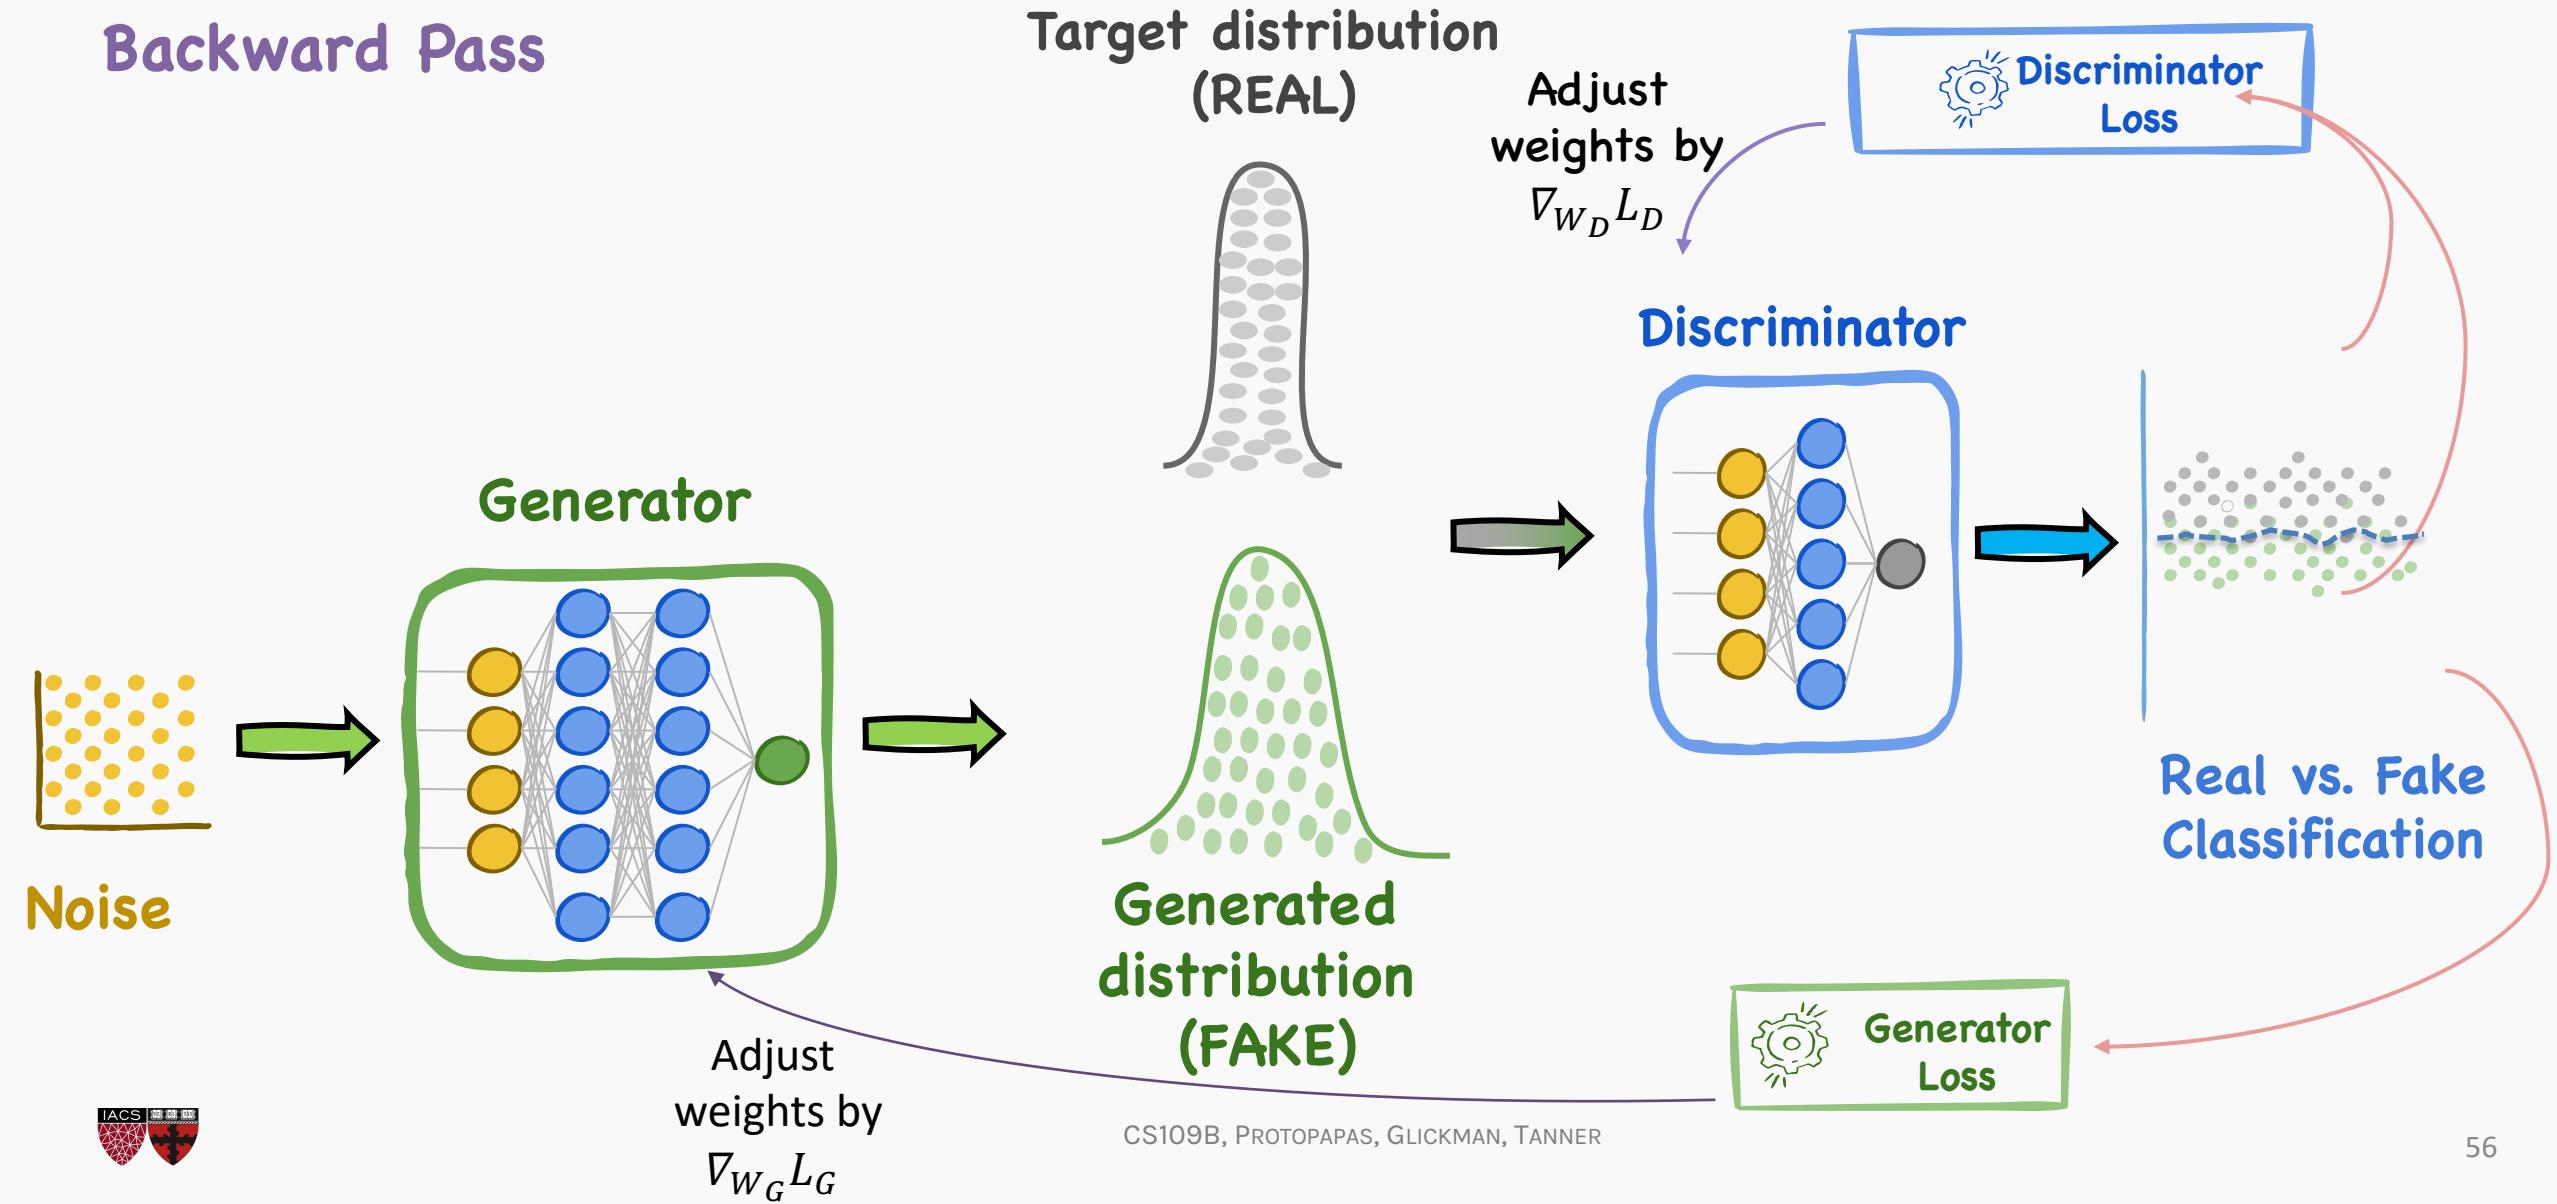
Training GANs

Backward Pass



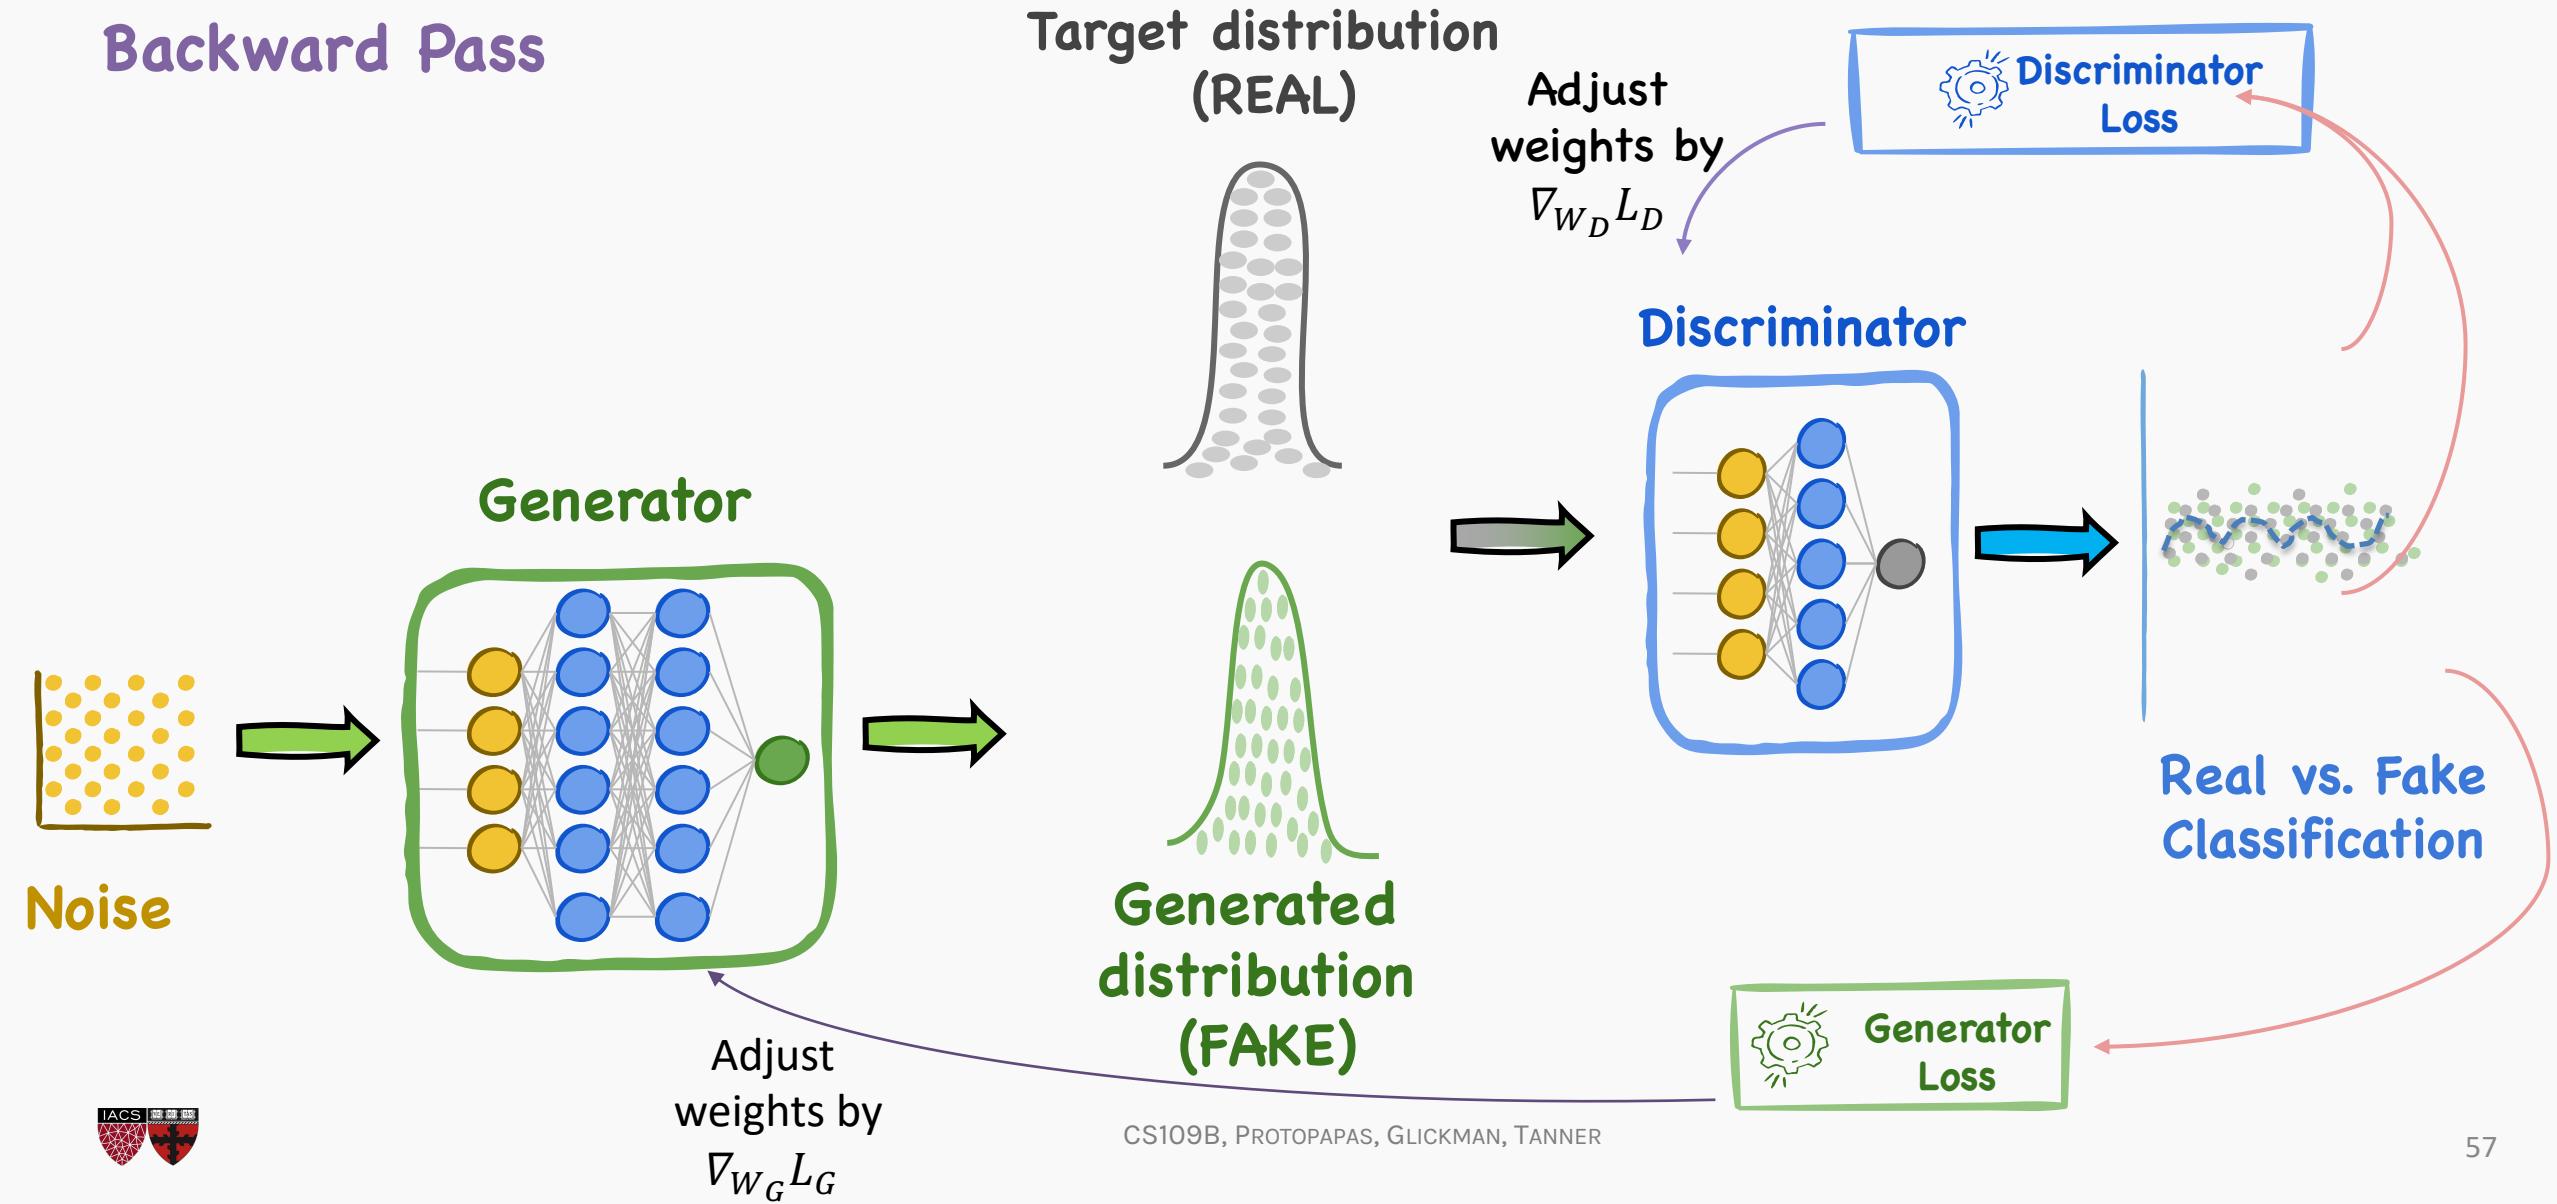
Training GANs

Backward Pass



Training GANs

Backward Pass



Generative Adversarial Networks

Training procedure

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

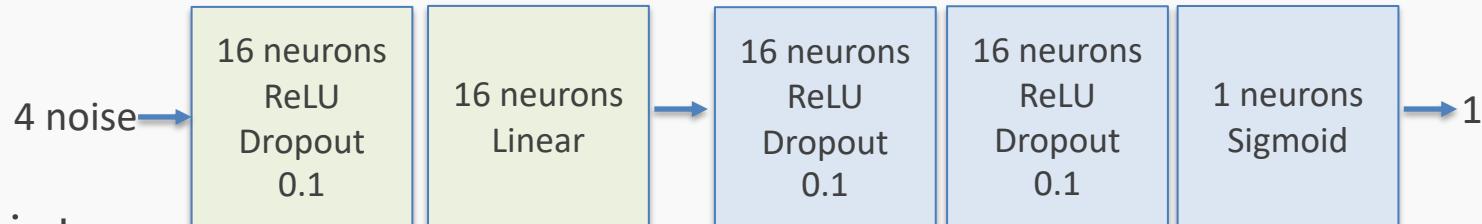
end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

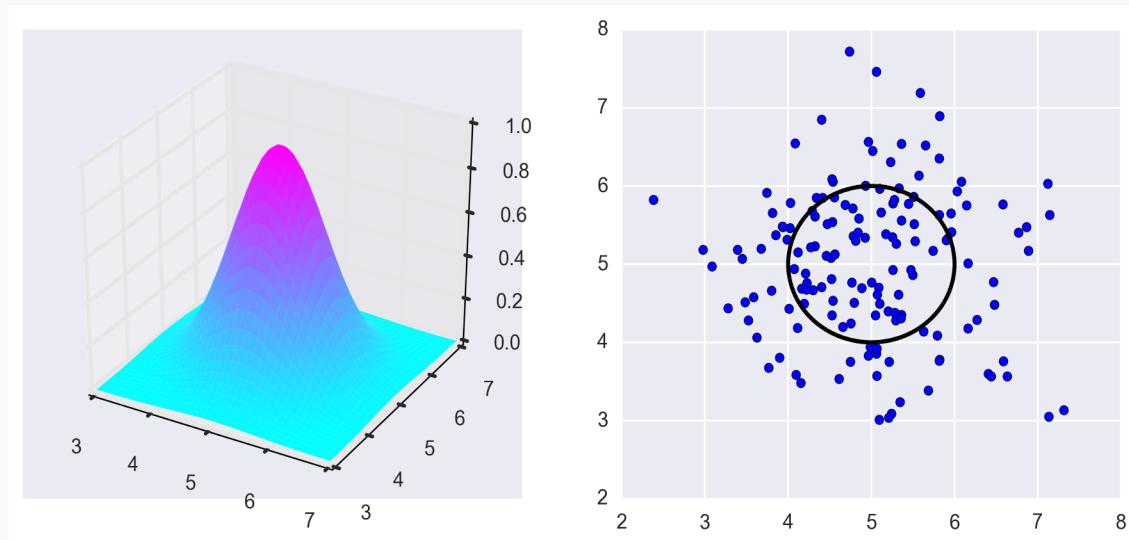
end for

Building GANS: Fully Connected Case



Let's build a FC simple GAN to generate points from a 2-dimensional Gaussian Distribution.

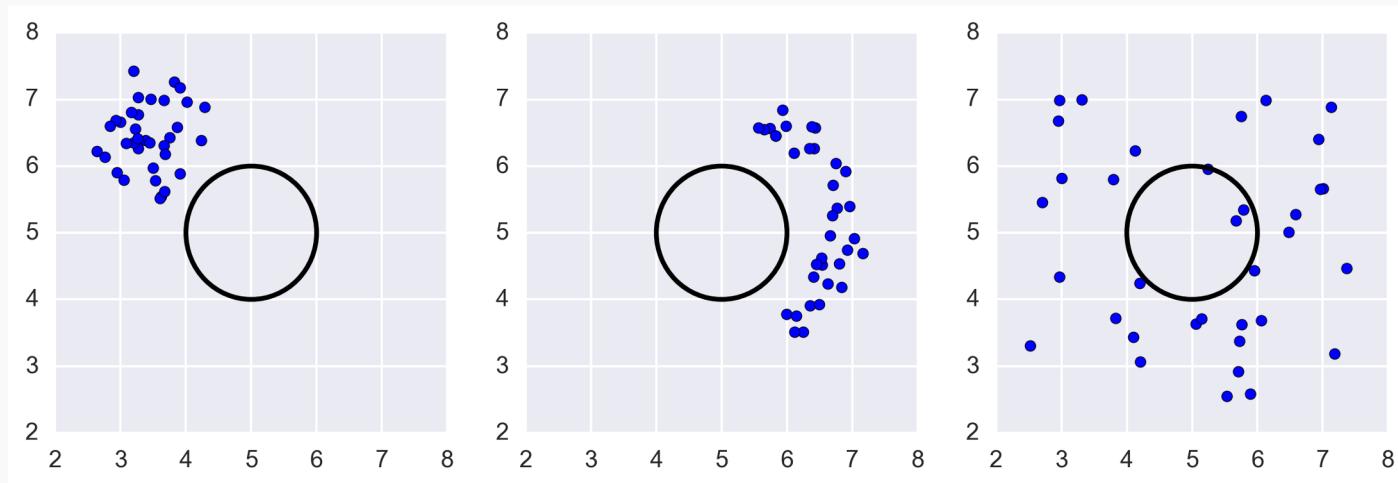
- **Generator**
 - Takes 4 random numbers
 - Generates a coordinate pair
- **Discriminator**
 - Takes an input point in the form of a coordinate pair
 - Determines whether the point is drawn from a specific 2-D Gaussian



Building GANS: Fully Connected Case

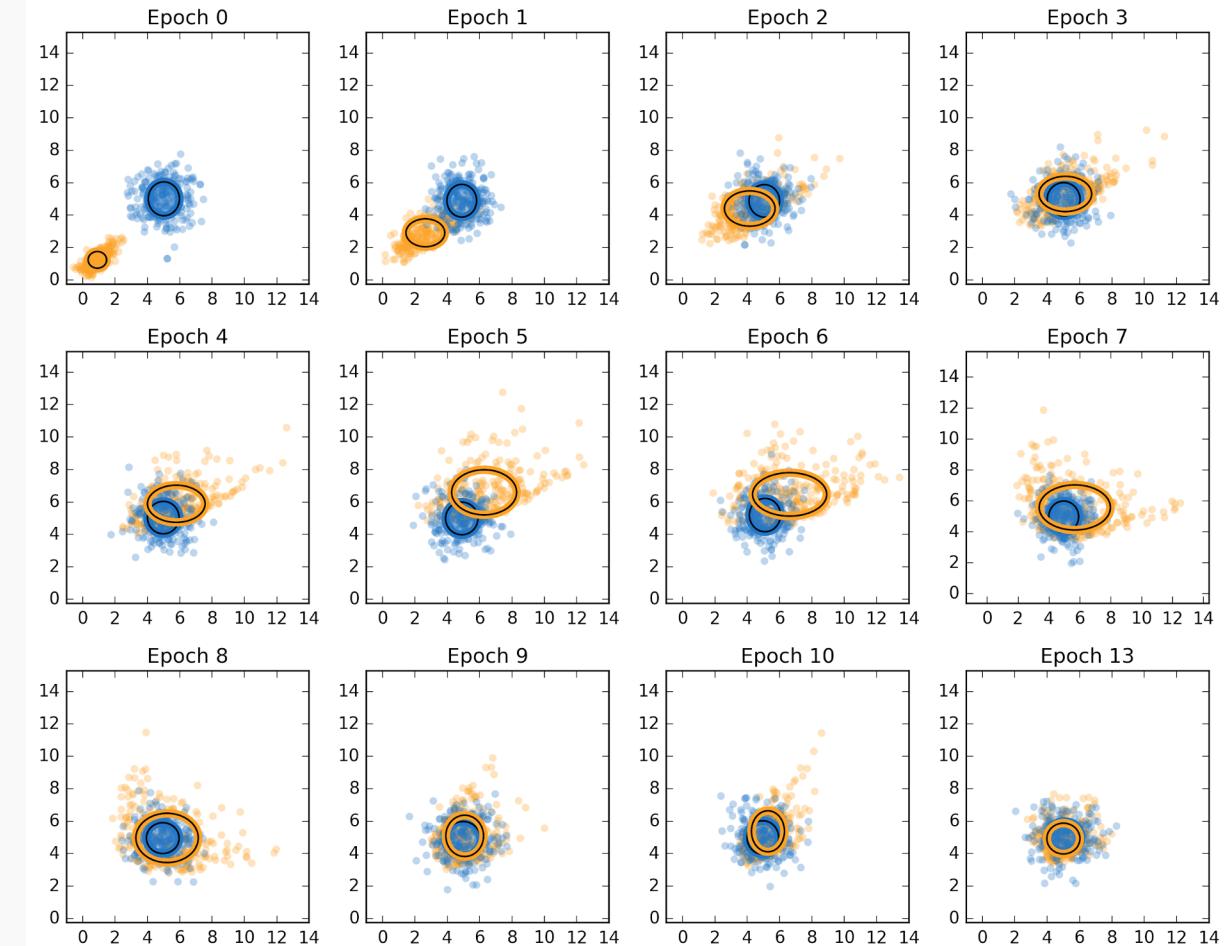
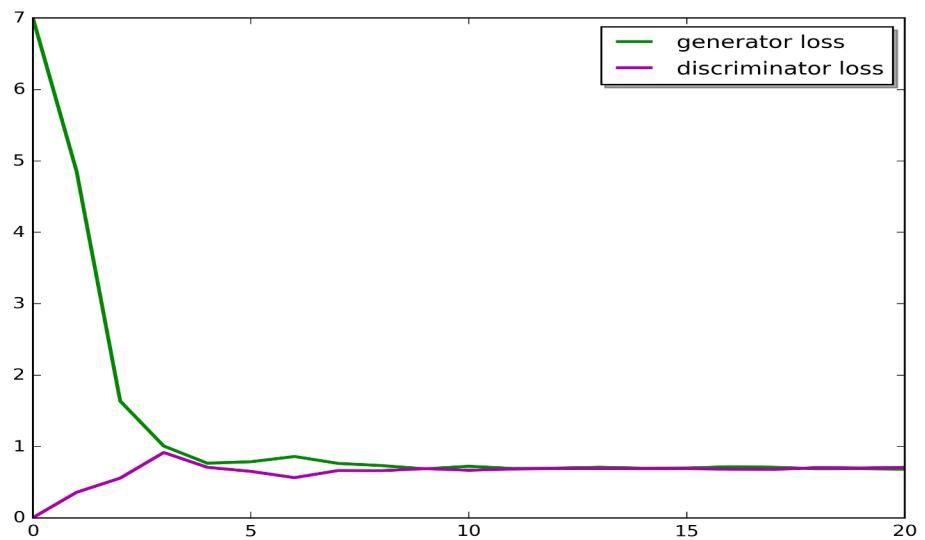
Train the Networks based on their ability to generate/discriminate batches of points drawn from the distribution.

Are these batches of points drawn from the right distribution?



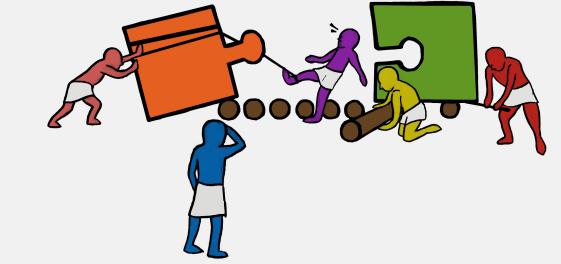
Building GANS: Fully Connected Case

As the generator and discriminator loss converges, the batch of points generated by the generator (in the yellow) approaches the real batch of points (in the blue)



Exercise:

In this exercise, we are going to generate 1-D Gaussian distribution from a n-D uniform distribution. This is a toy exercise in order to understand the ability of GANs (generators) to generate arbitrary distributions from random noise.

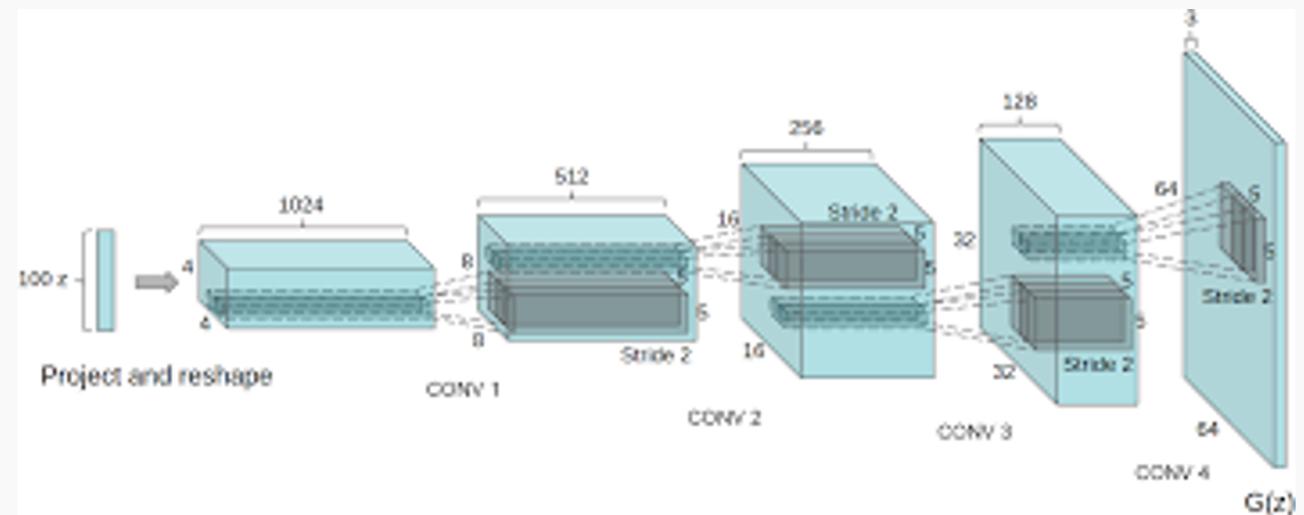


Deep Convolutional GAN: DCGAN

Deep Convolutional GAN (DCGAN)

-- Alex Radford et al. 2016

- Eliminate fully connected layers.
- Max Pooling BAD! Replace all max pooling with convolutional stride
- Use transposed convolution for upsampling or simply upsampling.
- Use Batch normalization



Building GANS: DCGAN

DCGAN on MNIST

Generated digits

6	0	3	1	3	5	3	6	1	1
1	5	3	0	5	6	5	9	3	0
8	1	7	9	6	7	2	5	6	6
3	2	0	4	2	1	4	0	7	7
0	1	7	5	0	9	1	0	0	0
7	3	8	9	1	9	9	5	3	2
7	8	3	2	2	8	1	4	5	9
1	5	4	5	0	7	9	6	8	5
3	0	1	8	4	3	9	1	0	0
9	8	6	7	6	3	1	8	0	0

Evolution of GANs

5 Years of Improvement in Artificially Generated Faces



https://twitter.com/goodfellow_ian/status/969776035649675265?lang=en



Ian Goodfellow

@goodfellow_ian

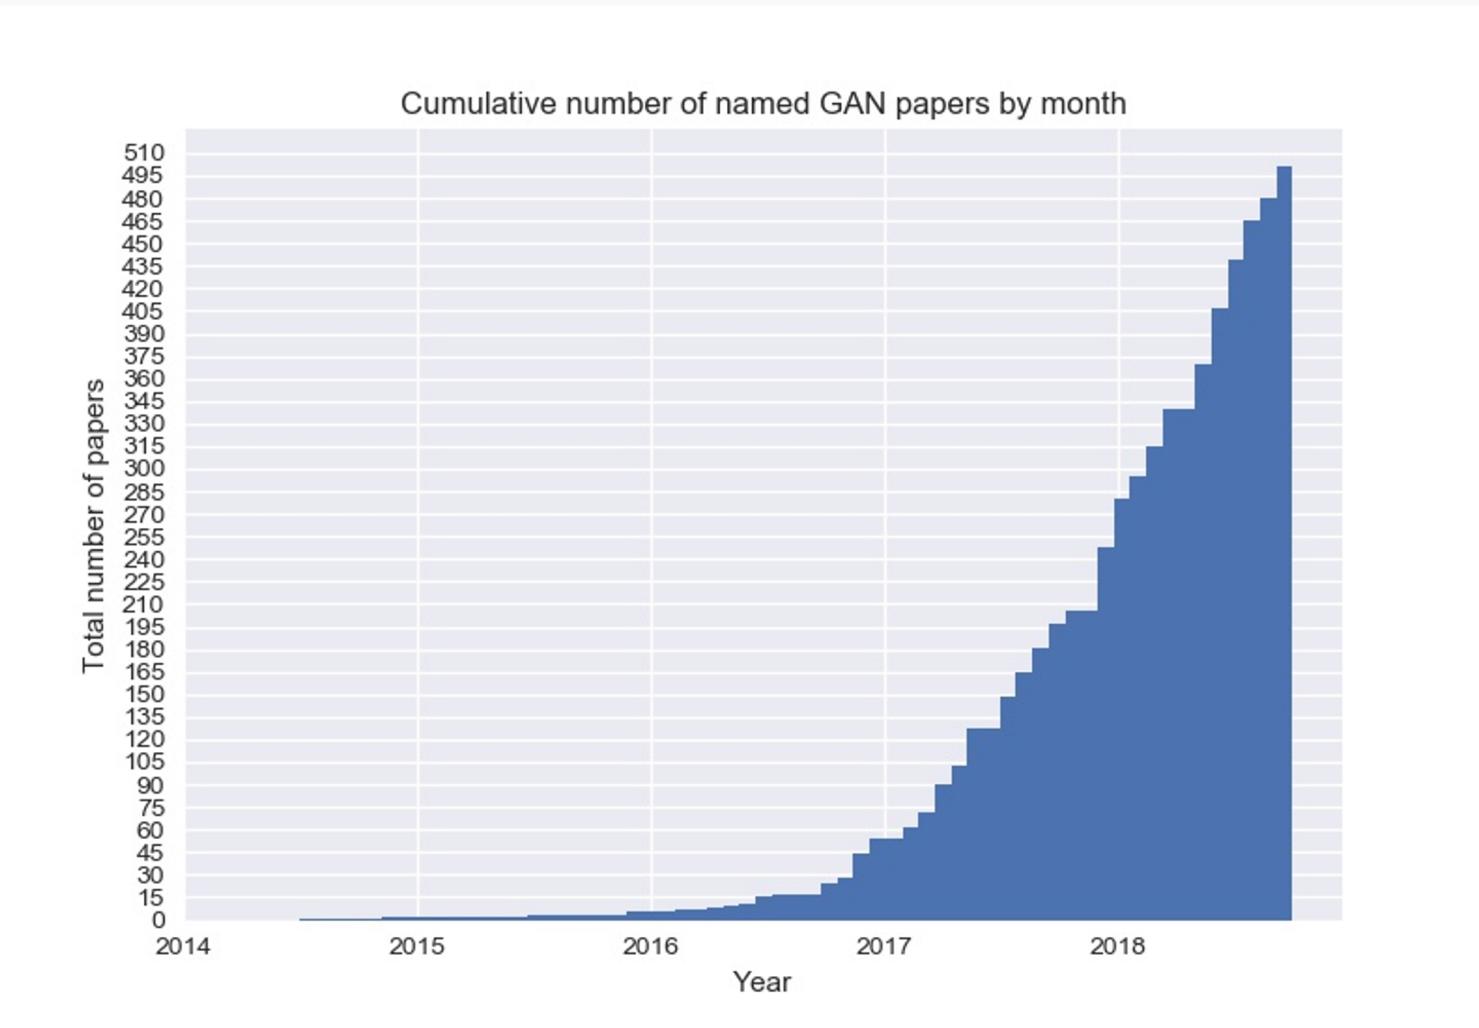
Follow



One of my favorite samples from the Progressive GANs paper is this one from the "cat" category. Apparently some of the cat training photos were memes with text. The GAN doesn't know what text is so it has made up new text-like imagery in the right place for a meme caption.



Evolution of GANs

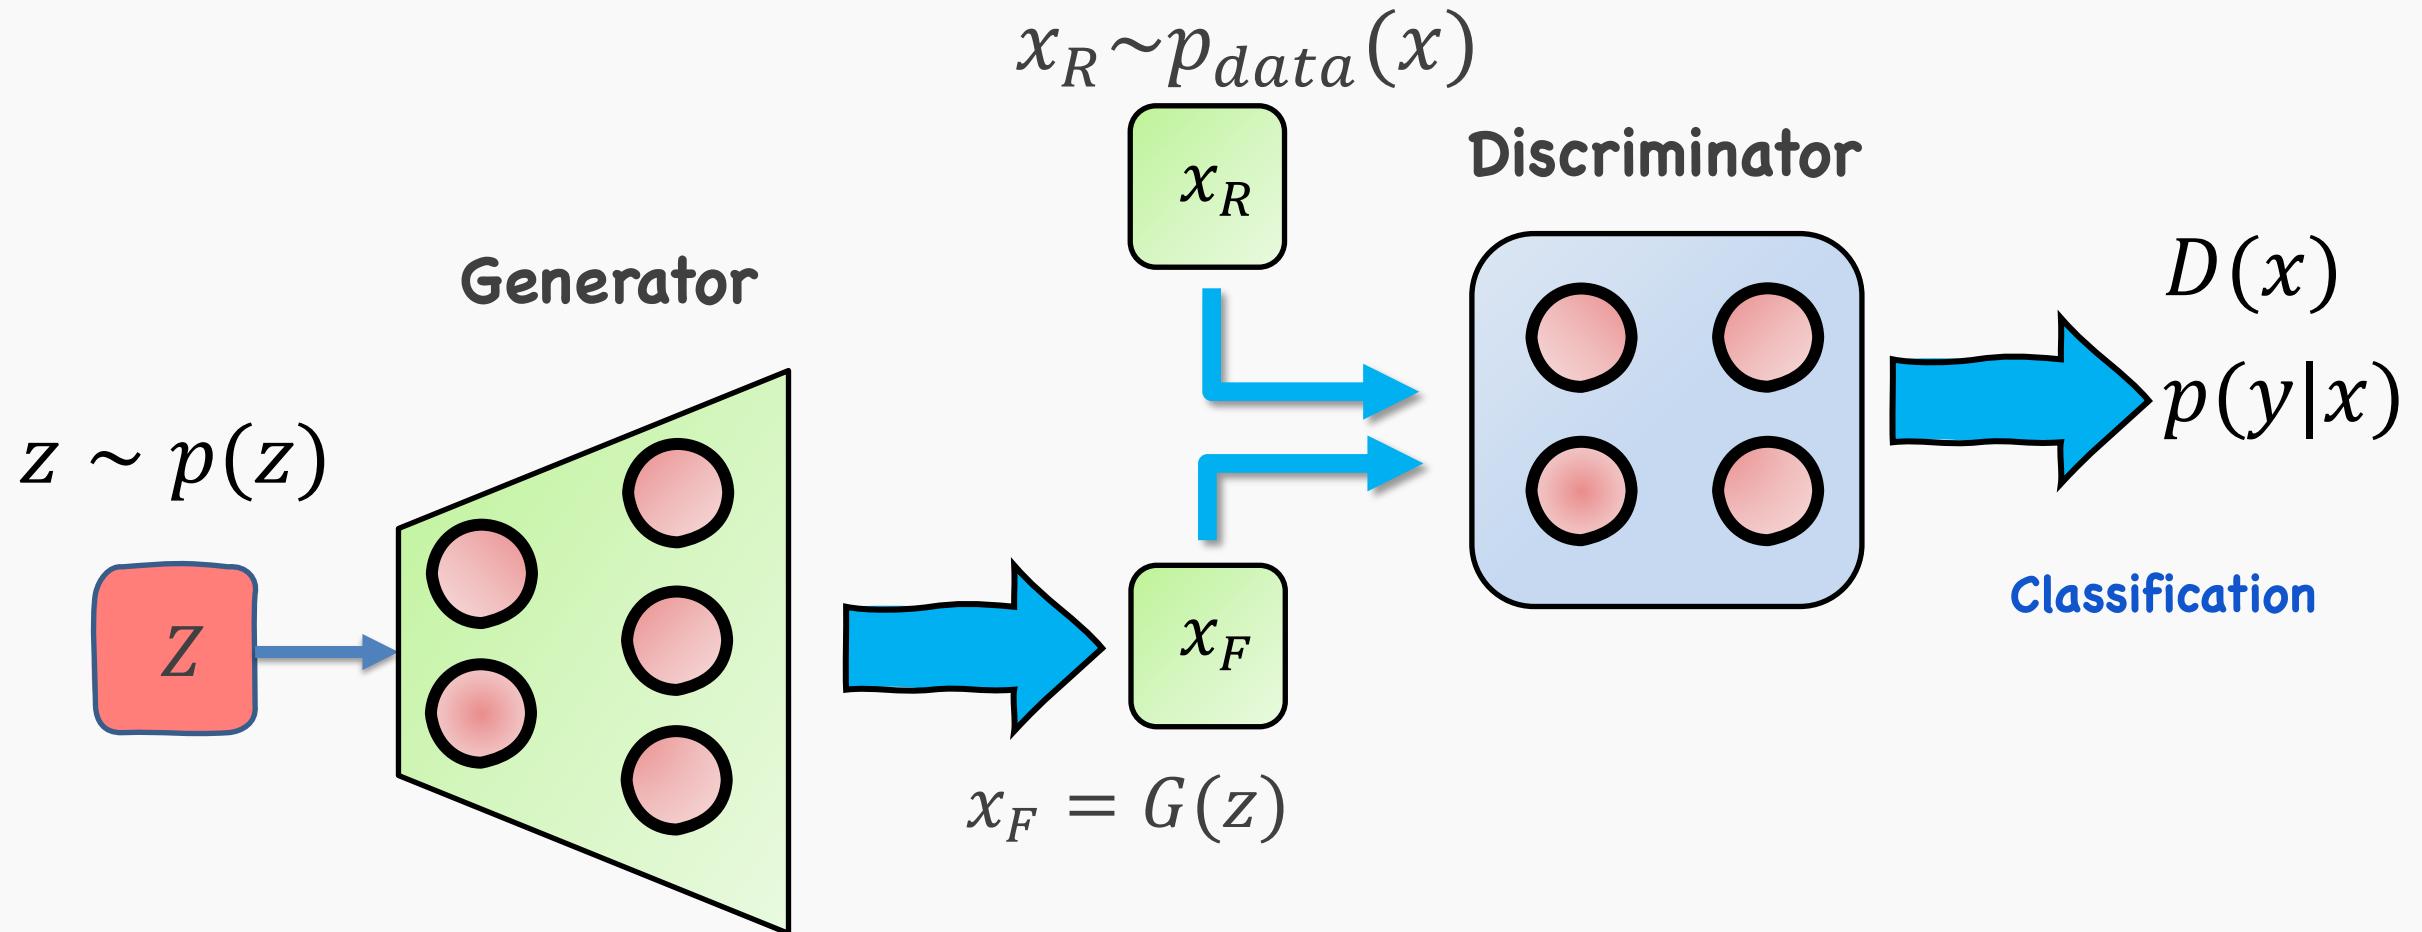


Outline

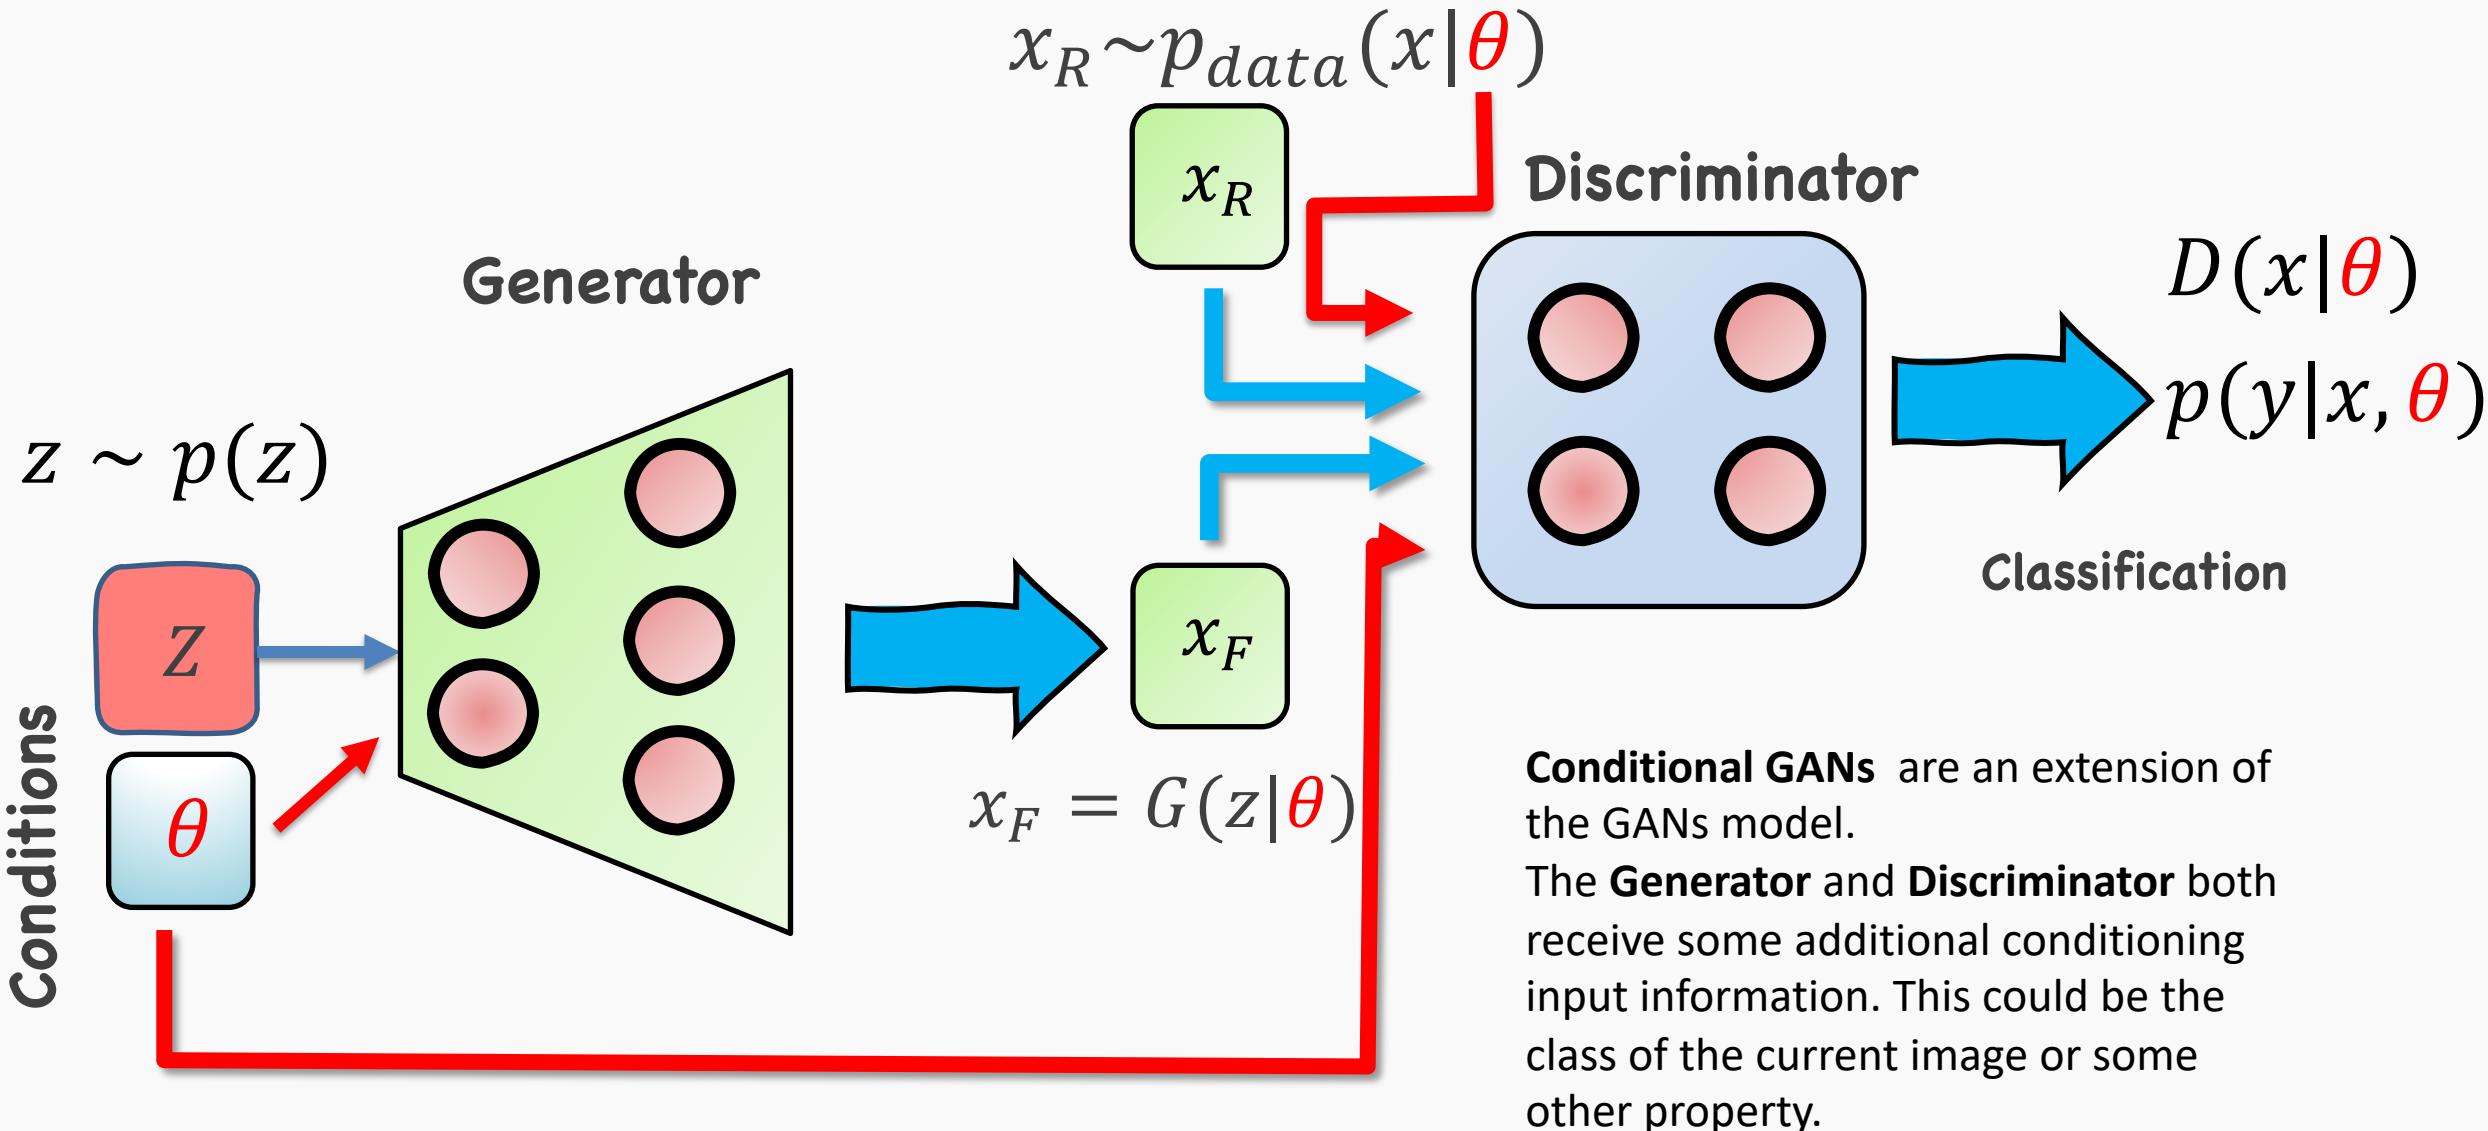
- Generative Modeling Motivation
- High Level Formalism
- Architecture
- Mathematics
- Training GANS
- **Conditional GANS**



Vanilla Generative Adversarial Nets



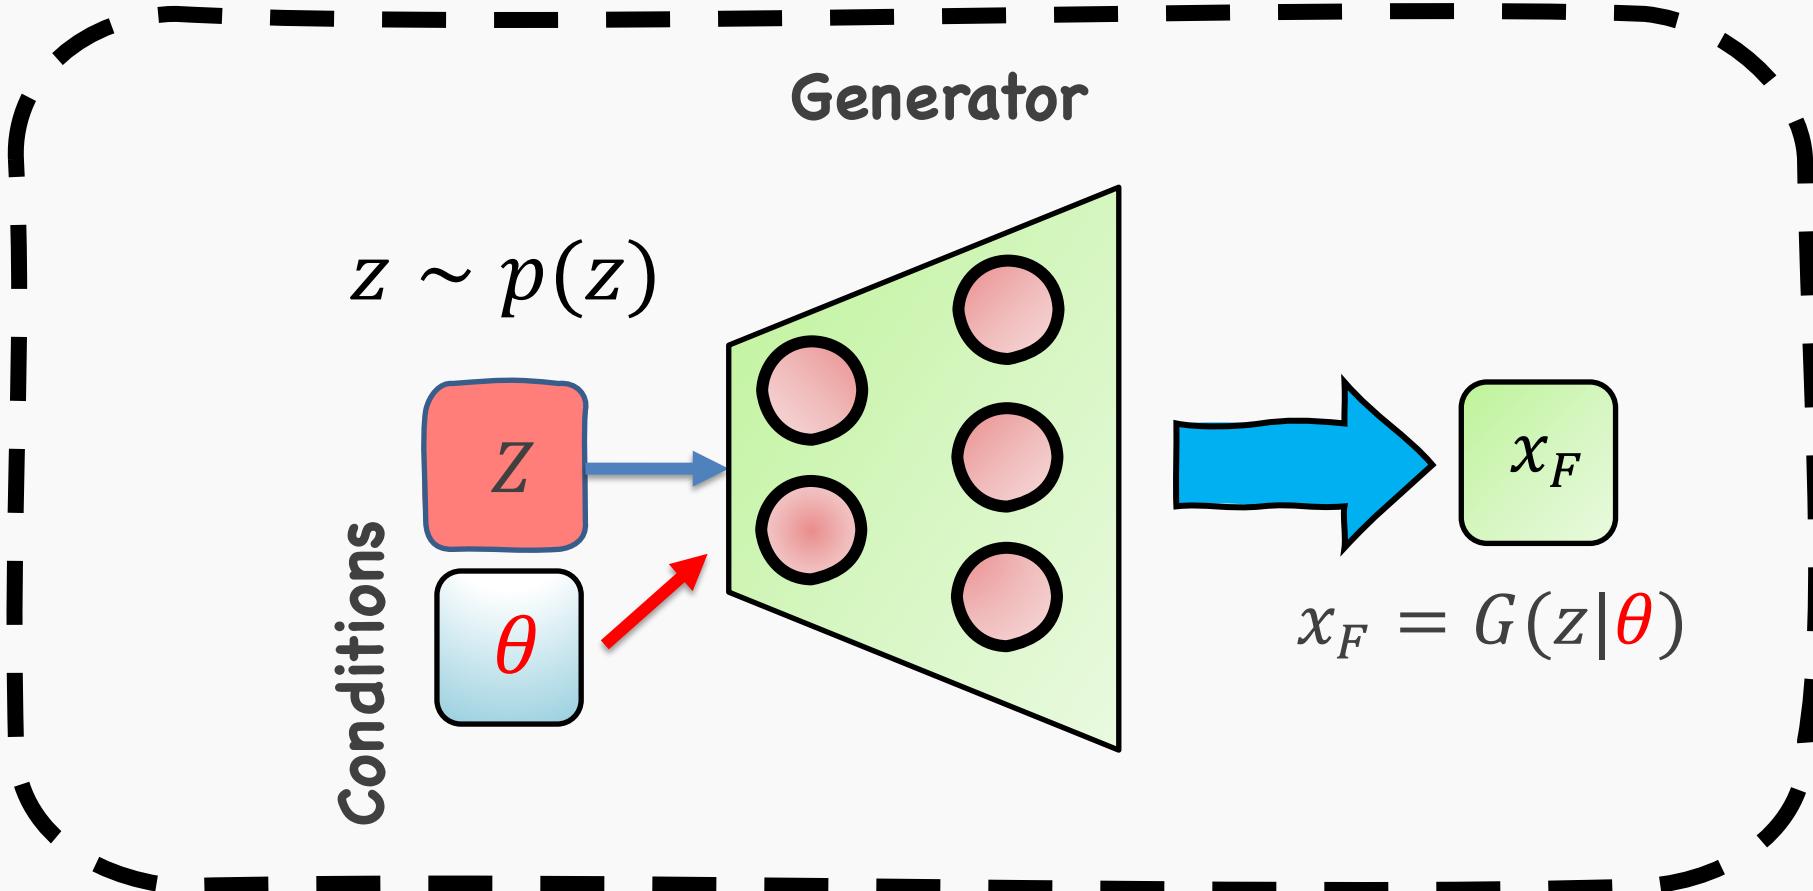
Conditional Generative Adversarial Nets



<https://arxiv.org/abs/1411.1784>



Conditional Generative Adversarial Nets



Next Lecture

- Challenges
 - Mode Collapse
- Evaluate GANs
- Auxiliary Classifier
- Wasserstein GAN
- Self-Attention Generative Adversarial Networks
- Large Scale GAN Training for High Fidelity Natural Image Synthesis